

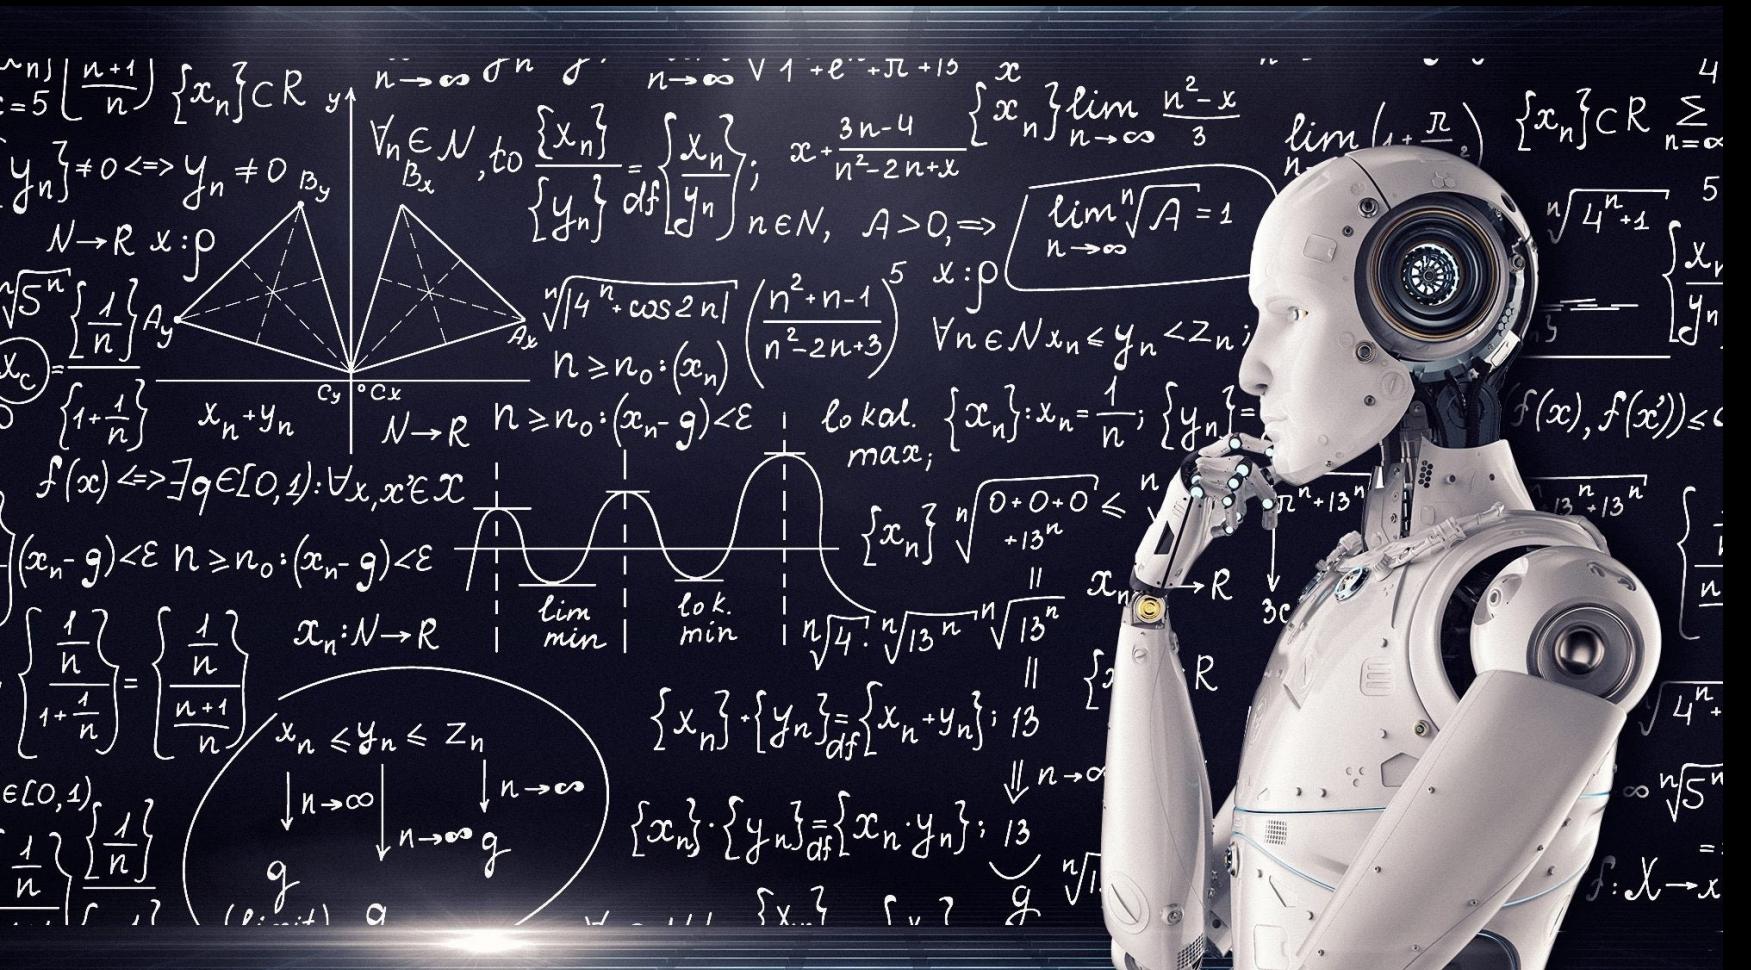
Design for
Artificial
Intelligence

ChatGPT 101

<https://github.com/iportilla/lowcode>

- Intros
- ChatGPT 101
- Business Examples
- Exercise





- ChatGPT 101
 - Business Examples
 - Exercise

Artificial Intelligence

Machine Learning

Deep Learning

Generative AI



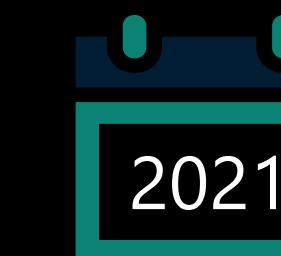
Artificial Intelligence
Intelligence demonstrated by machines



Machine Learning
Learn from data



Deep Learning
Model after the human brain (Neural Networks)



Generative AI
Create new written, visual, and auditory content

Generative AI

The best thing about AI is its ability to ...

learn	4.5%
predict	3.5%
make	3.2%
understand	3.1%
do	2.9%



OpenAI

GTP-4

Why is the sky blue?

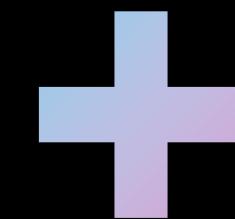
OpenAI GPTx Answer:

The sky blue because the Earth's atmosphere scatters sunlight in all directions and blue light is scattered more than the other colors because it travels as shorter, smaller waves

MSFT partnership with OpenAI



Ensure that artificial general intelligence (AGI) benefits humanity.



Empower every person and organization on the planet to achieve more

GPT-3

Generate and Understand Text

Codex

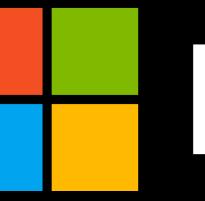
Generate and Understand Code

DALL·E

Generate images from text prompts



OpenAI



Microsoft

Generative AI

GPT-3

Prompt:

Write a tagline for an ice cream shop.

Response:

We serve up smiles with every scoop!

Codex

Prompt:

```
Table customers, columns =  
[CustomerId, FirstName,  
LastName, Company, Address,  
City, State, Country,  
PostalCode]
```

Create a SQL query for all customers in Texas named Jane
query =

Response:

```
SELECT *  
FROM customers  
WHERE State = 'TX' AND  
FirstName = 'Jane'
```

DALL·E

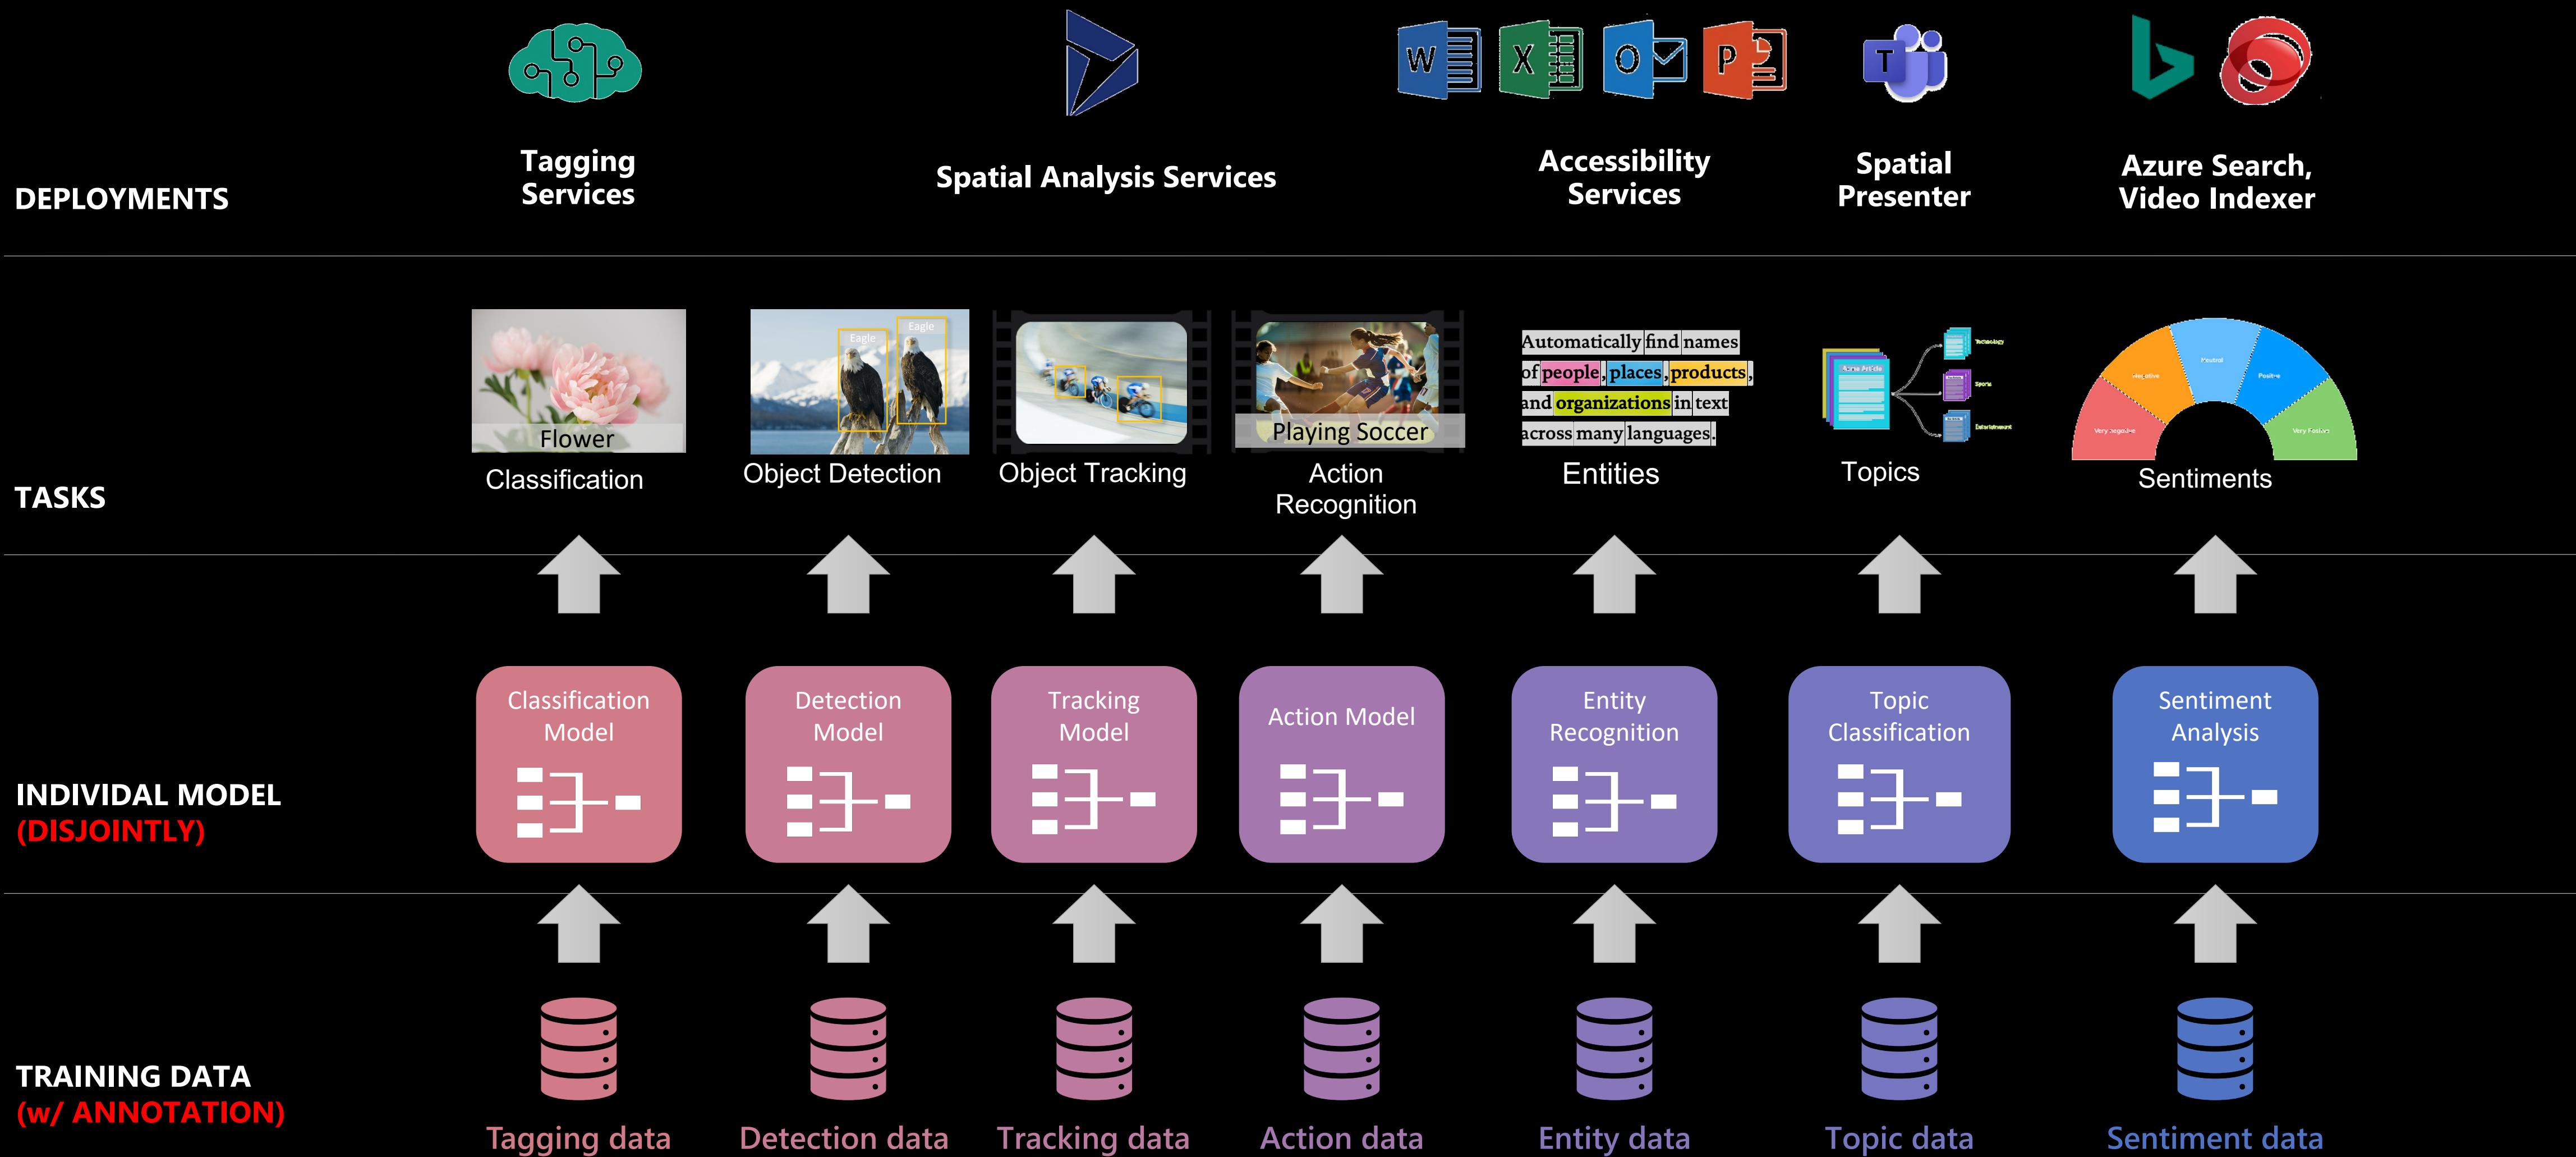
Prompt: A white Siamese cat

Response:

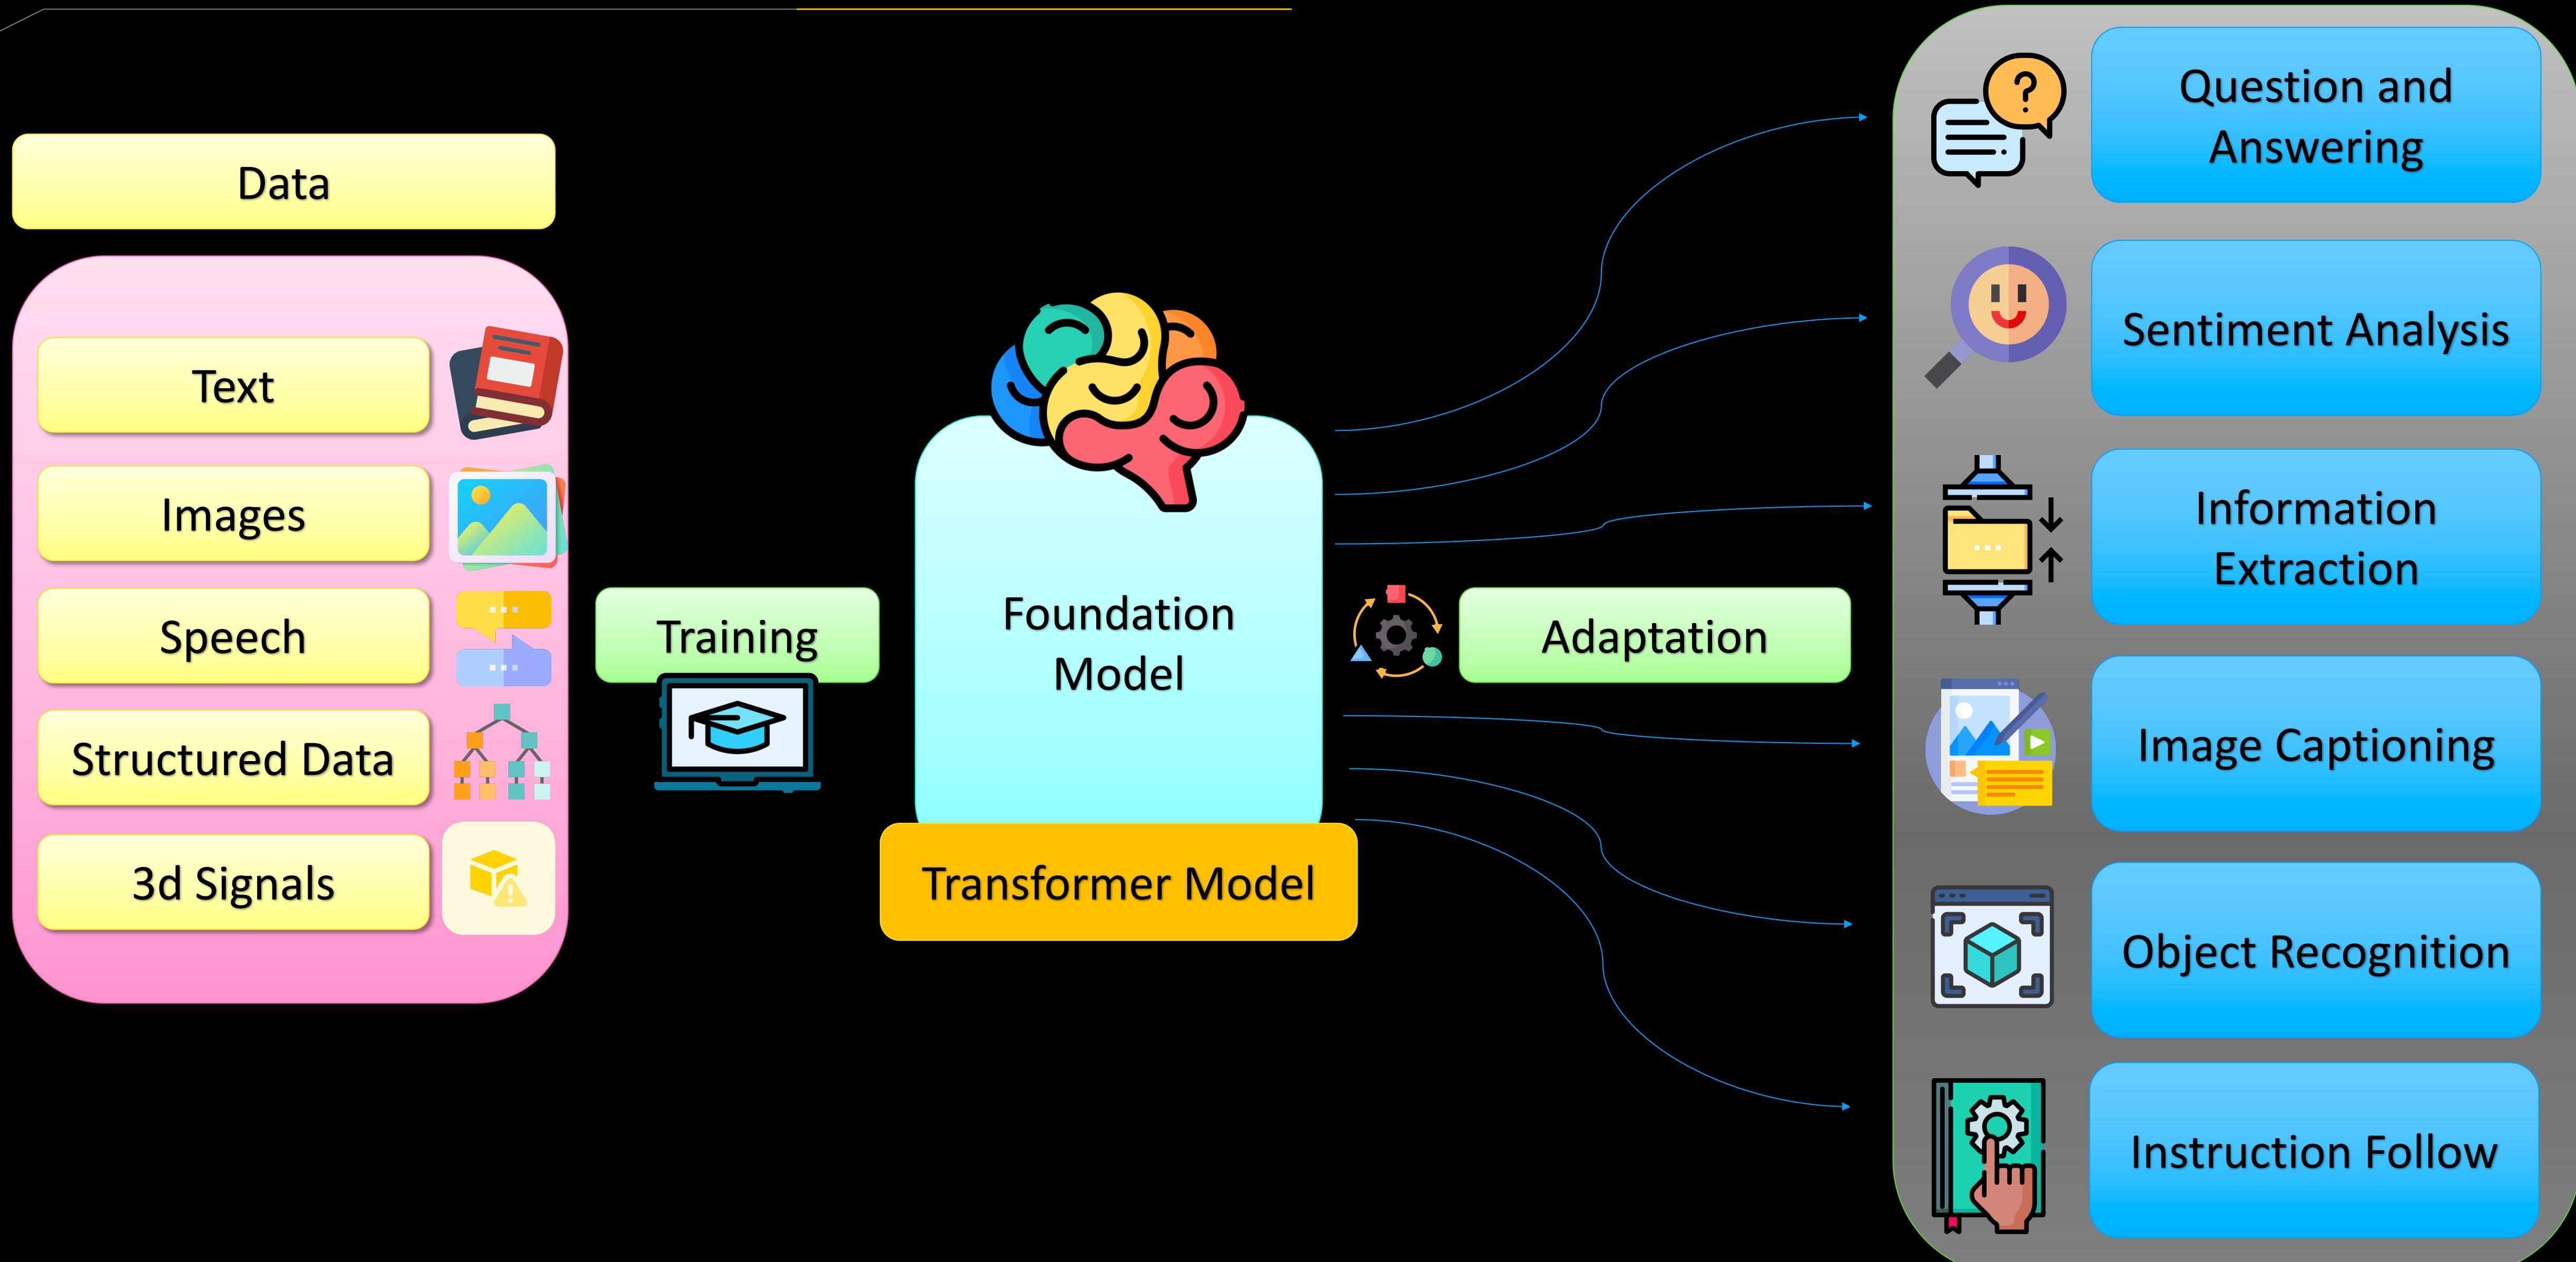


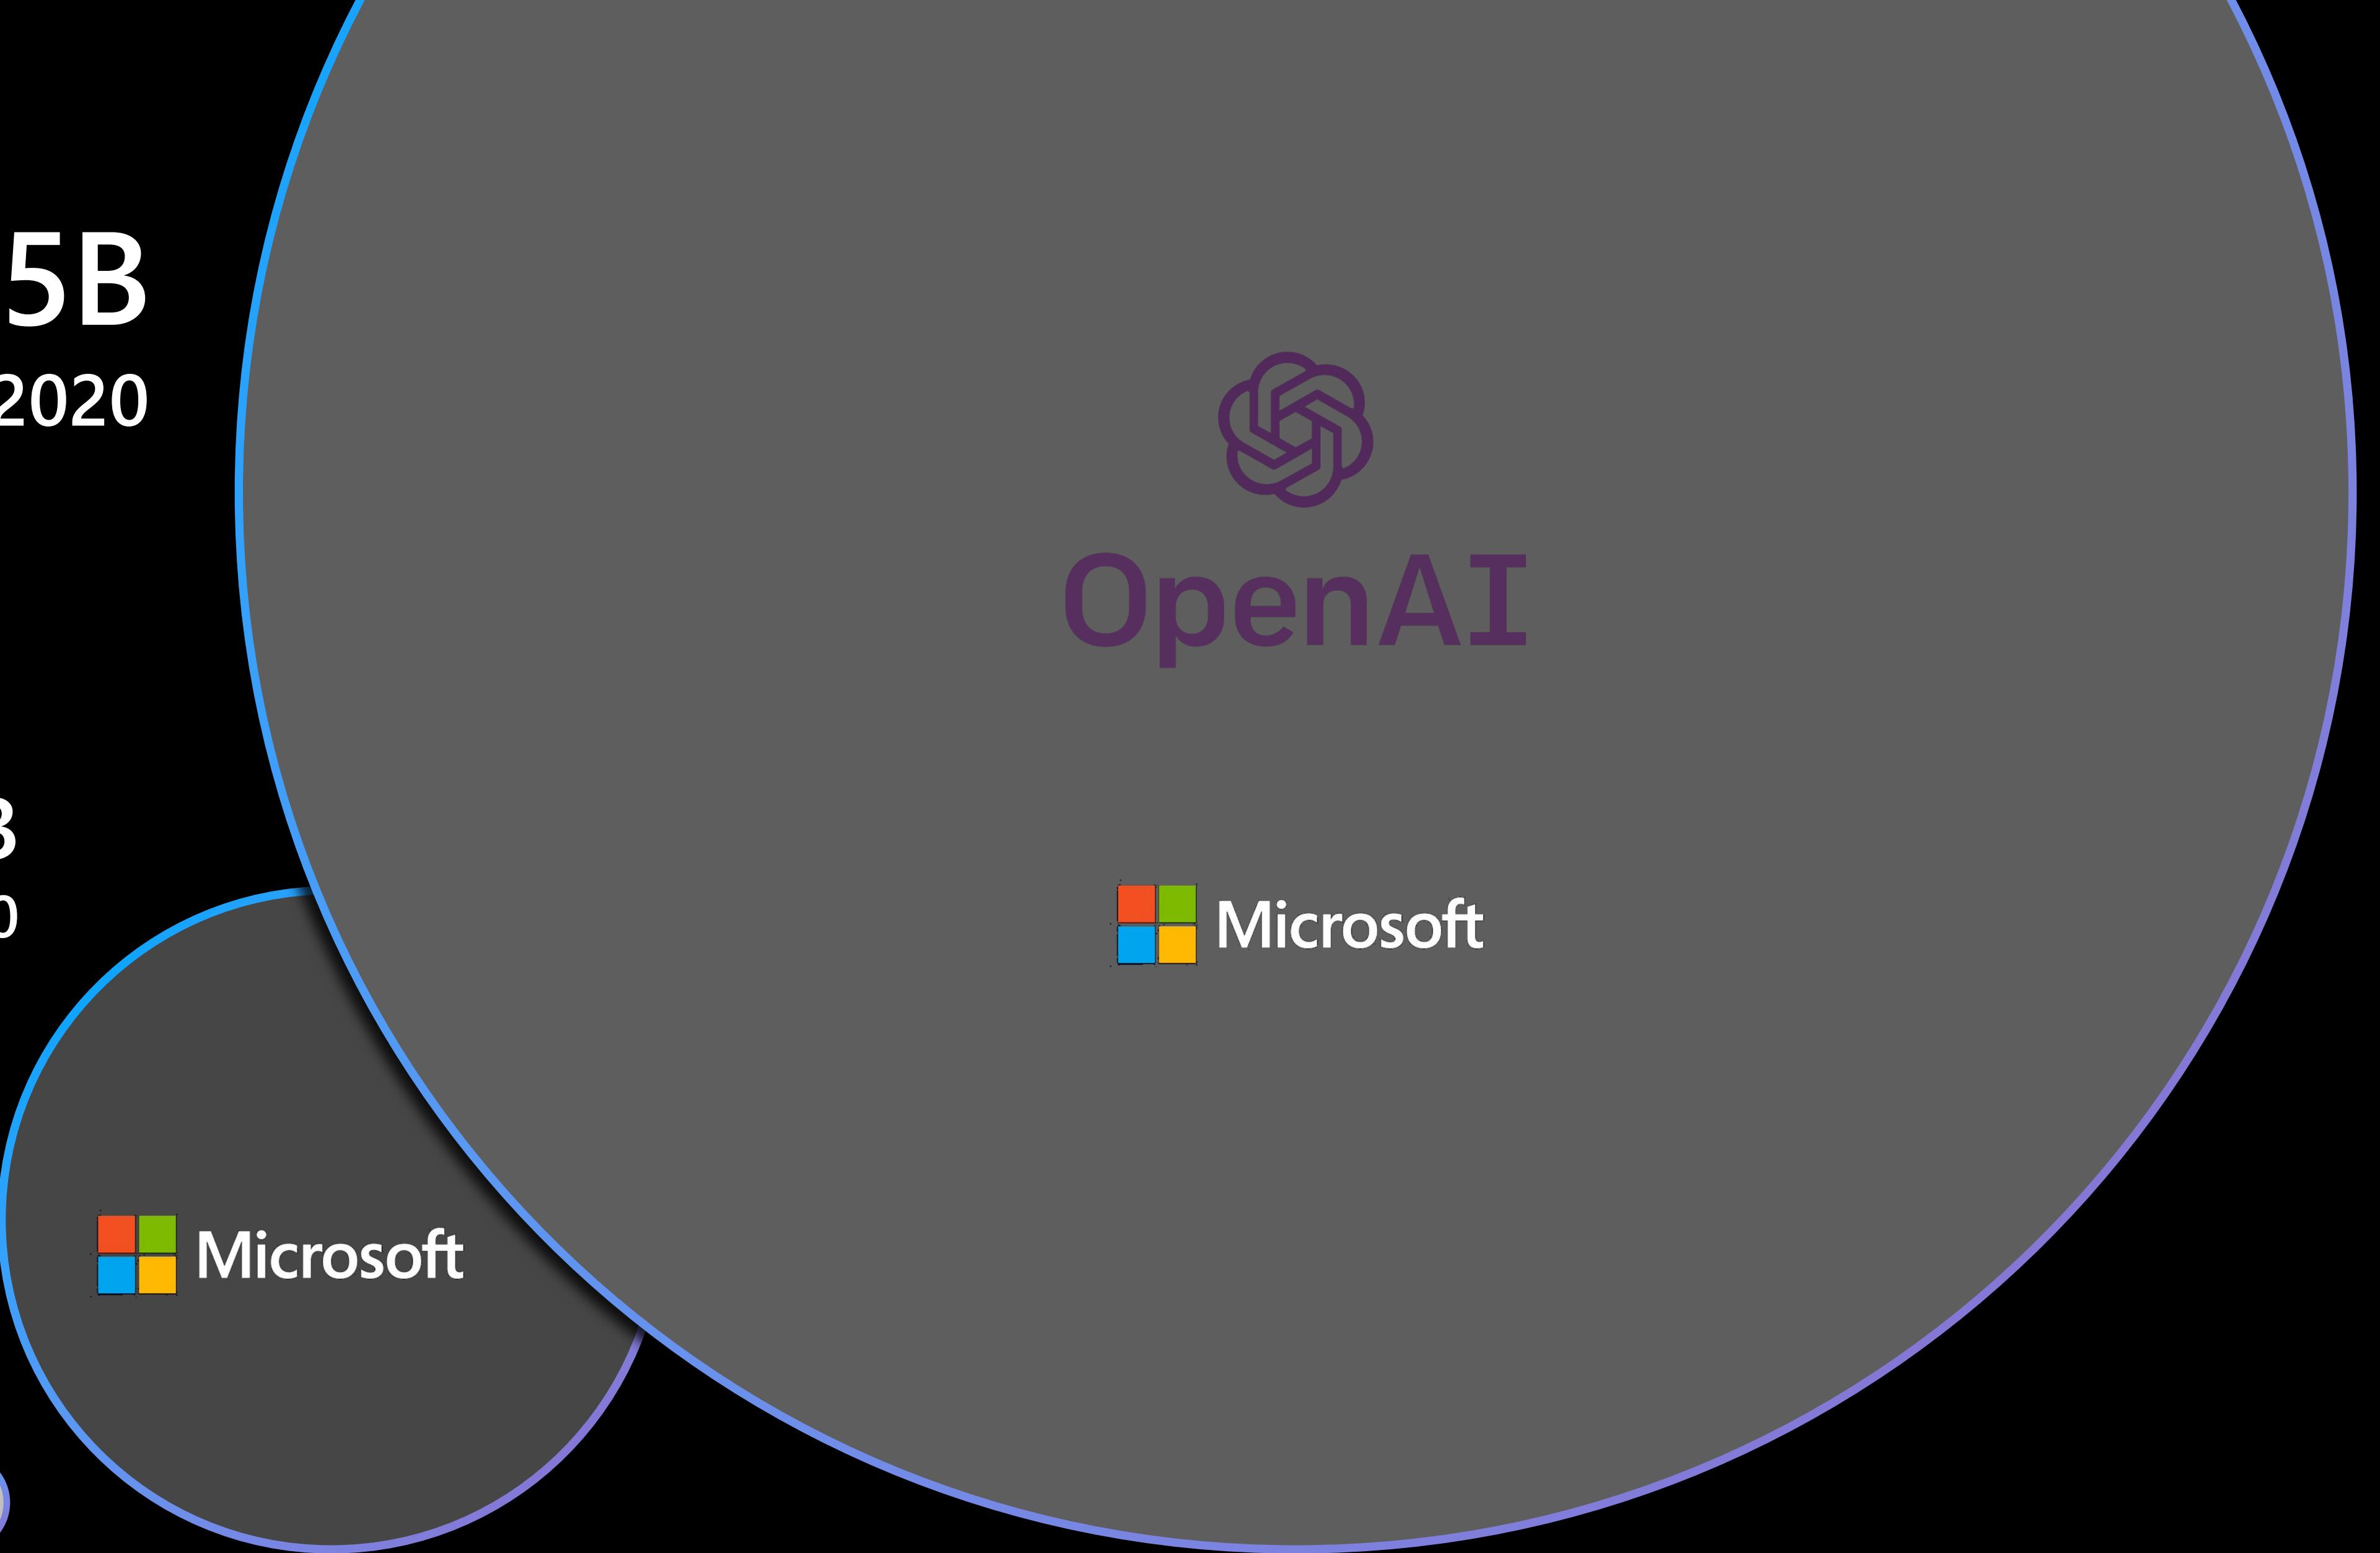
Traditional model development

High cost & slow deployment - Each service is trained disjointly



Foundation Models





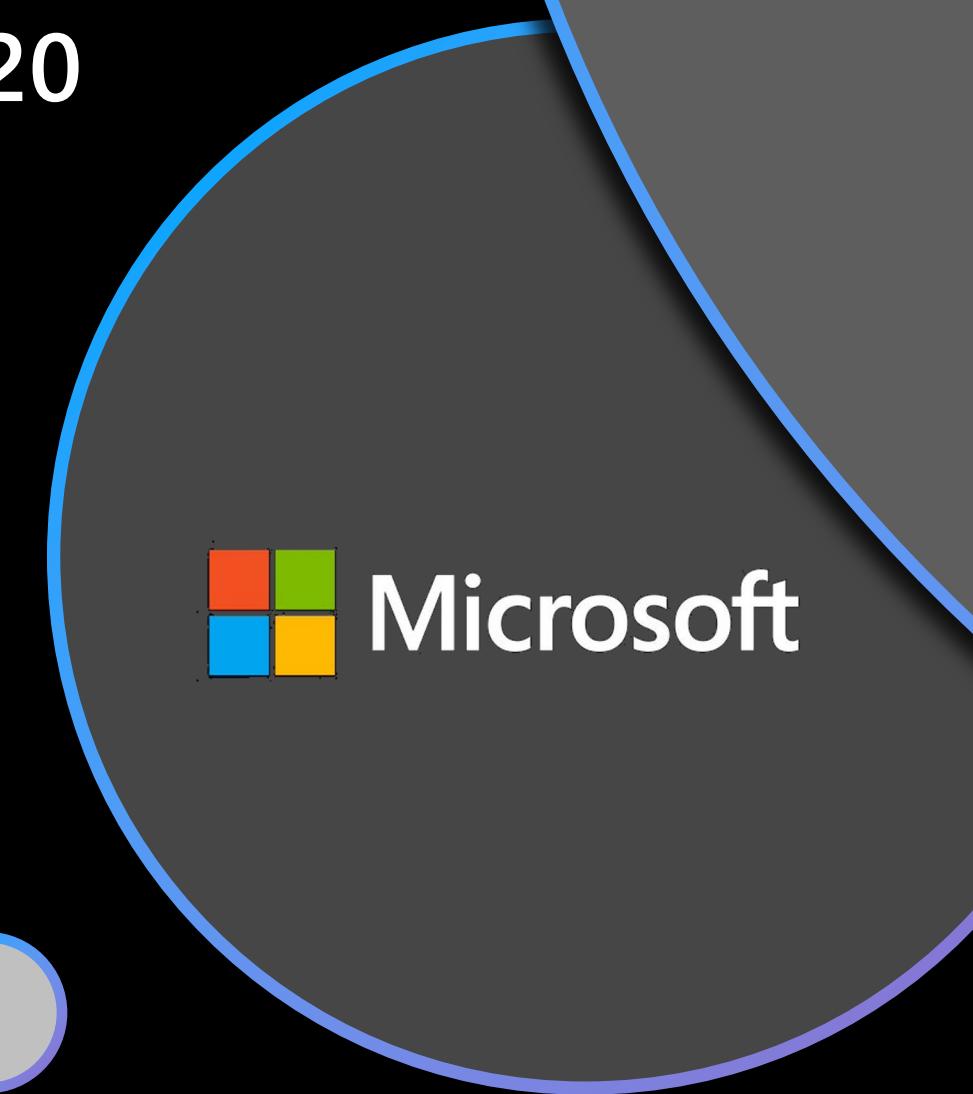
175B

GPT-3 | 2020



17B

Turing-NLG | 2020



340M

BERT-Large | 2018



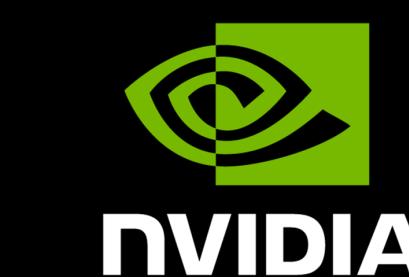
OpenAI



Microsoft



Microsoft



530B

Megatron-Turing
NLG | 2021

175B

GPT-3 | 2020

17B

Turing-NLG | 2020

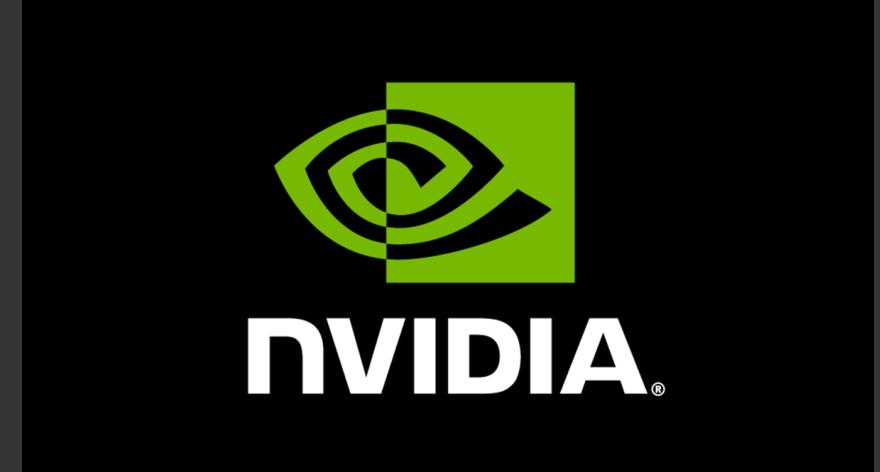
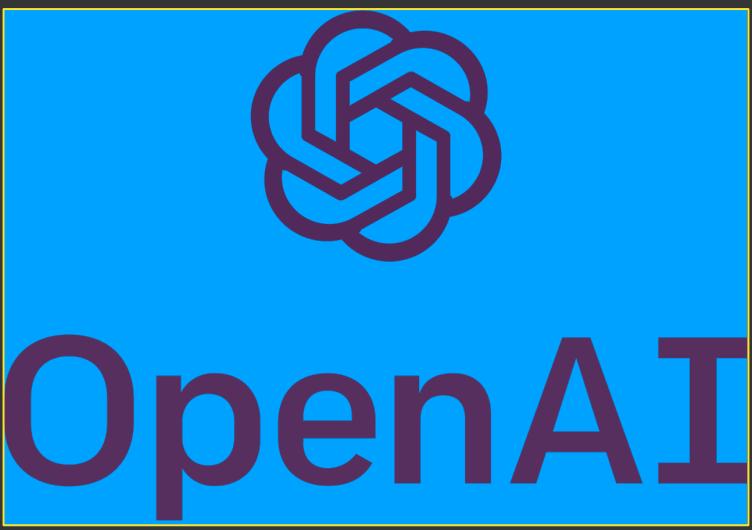
340M

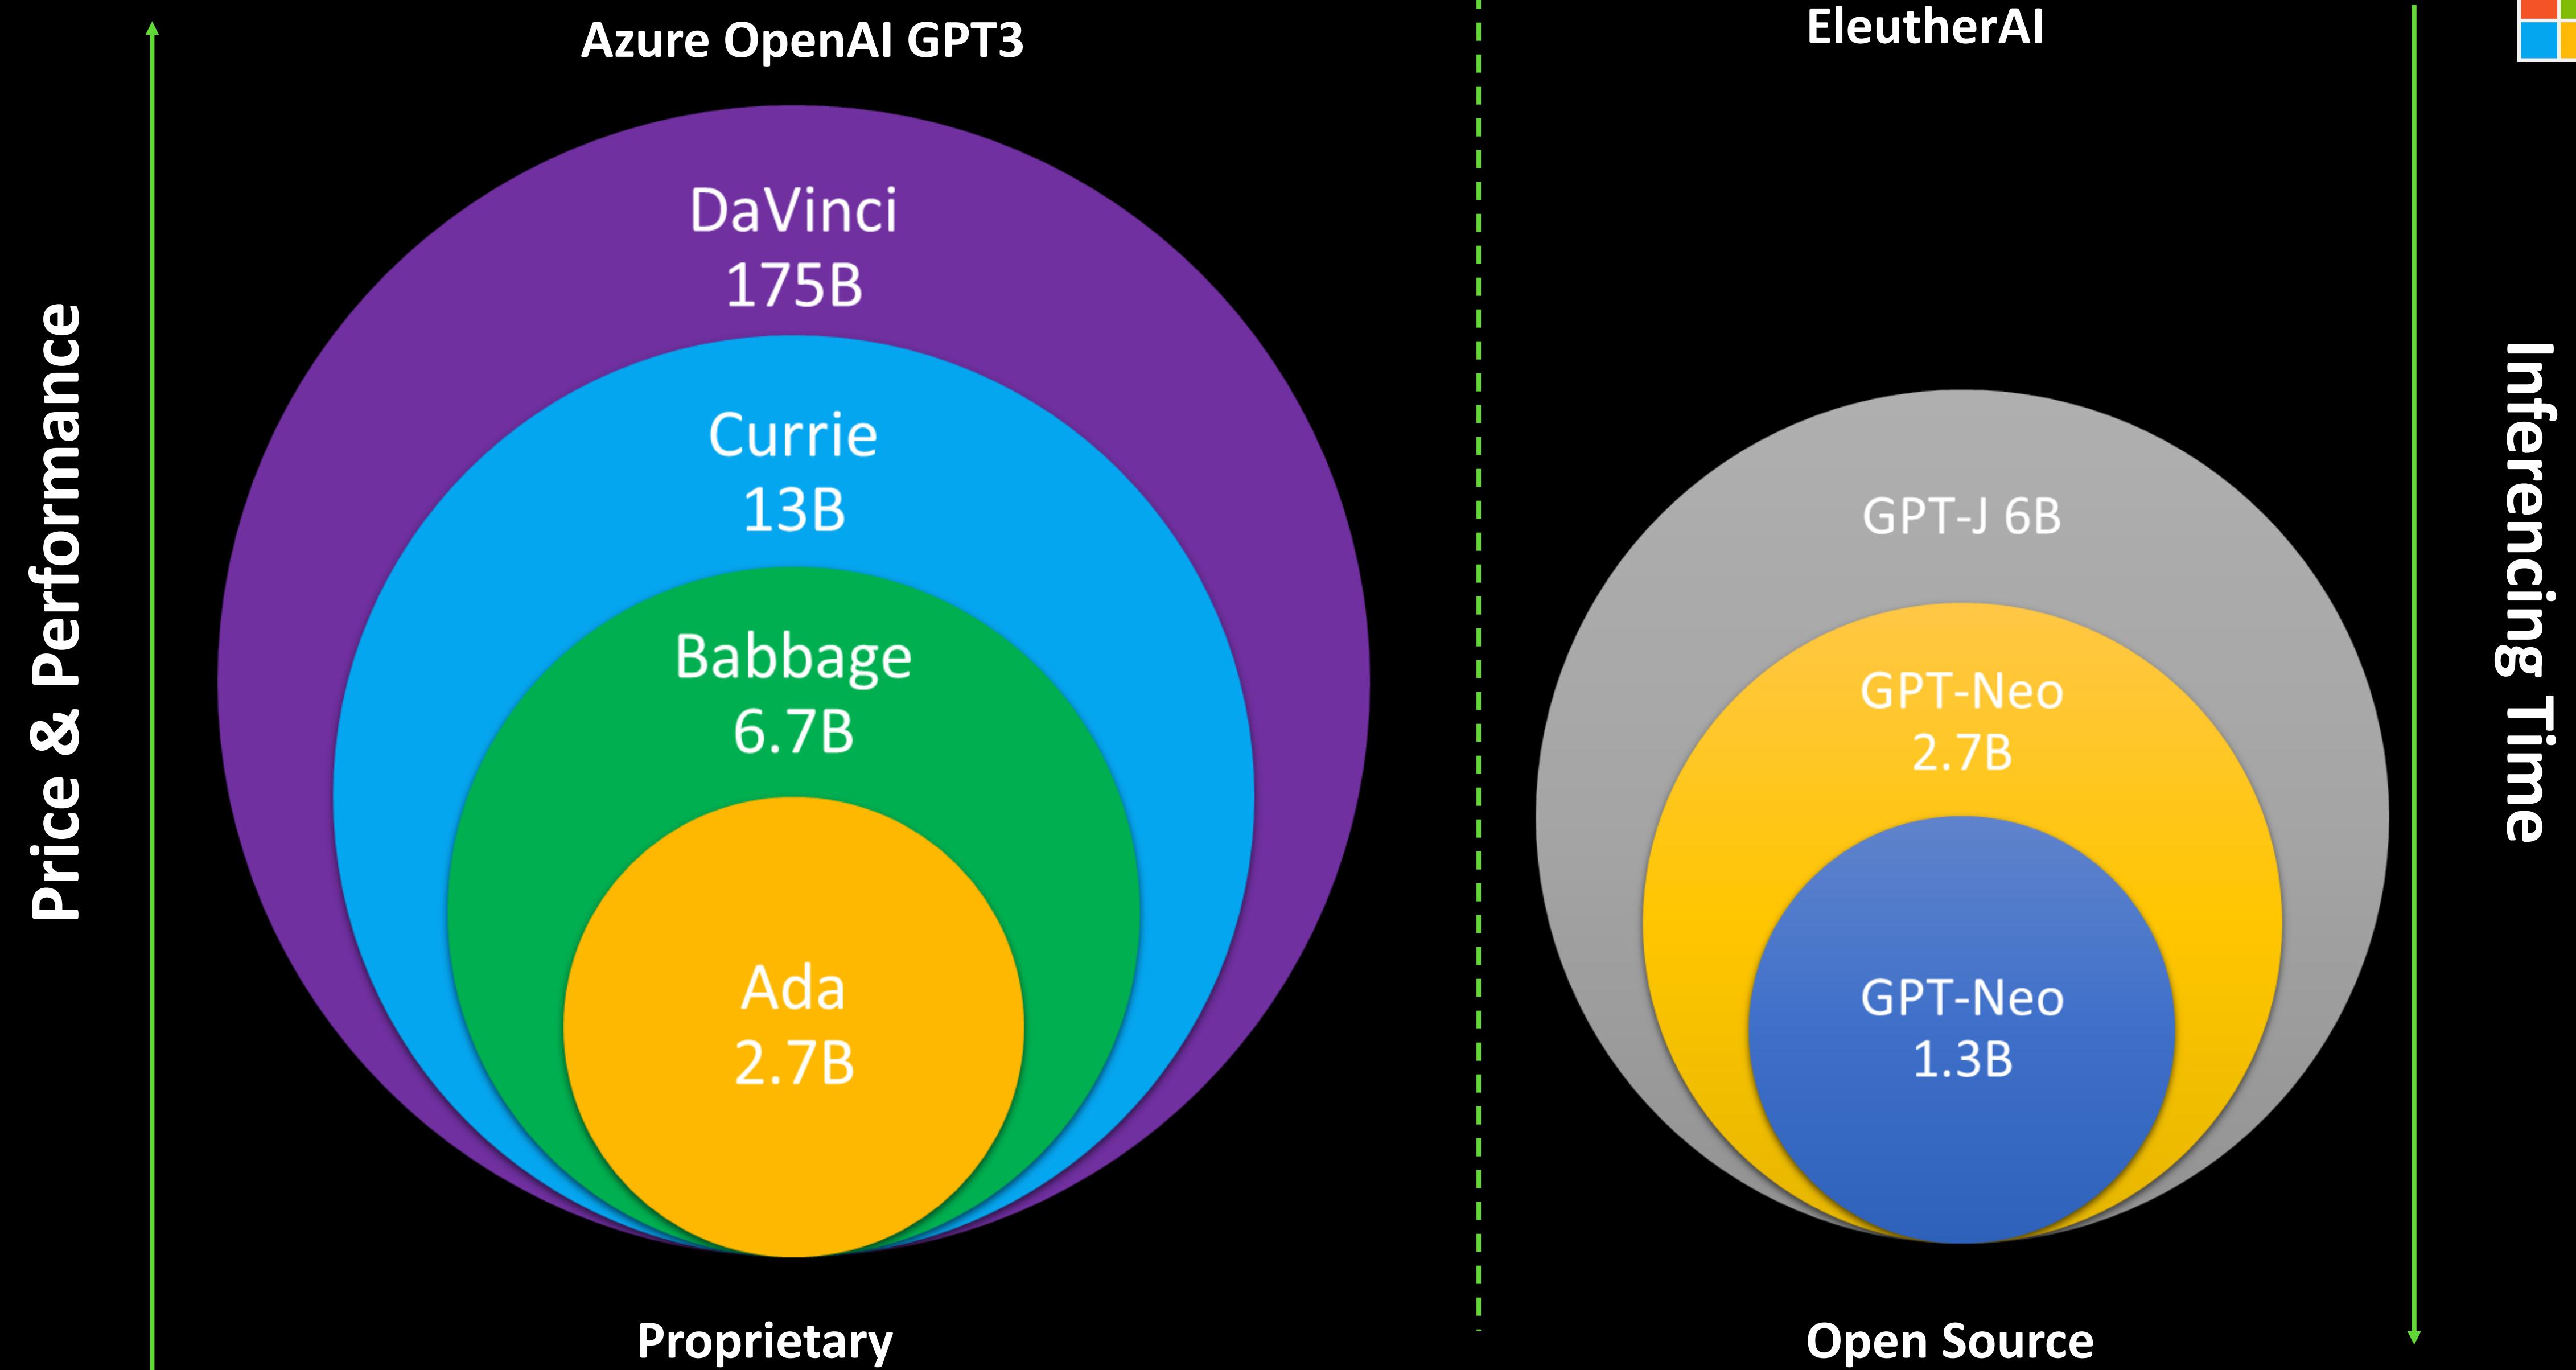
BERT-Large | 2018

530B
GPT-3 | 2020

1.8 T

ChatGPT 4.0
ChatGTP Turbo



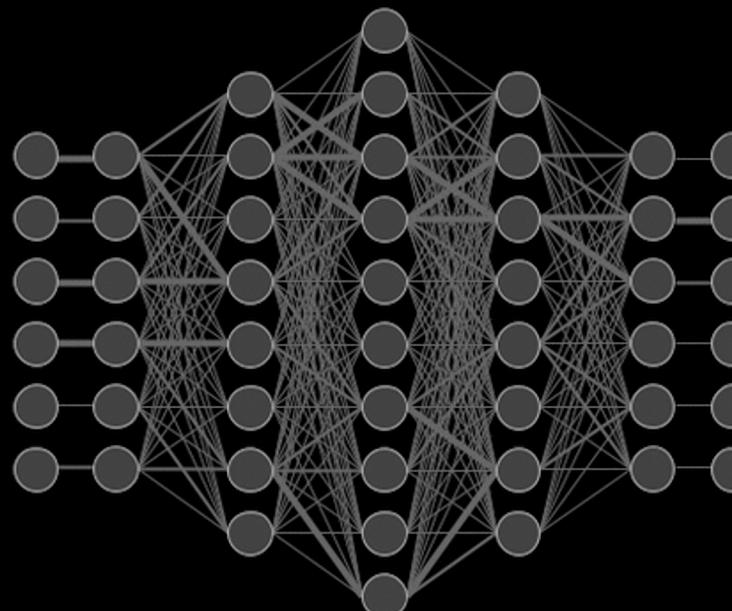


- Intros
- ChatGPT 101
- Business Examples
- Exercise

Model use out of the box - Prompting

Auto-Regressive
(Predicts next word)

Decoder

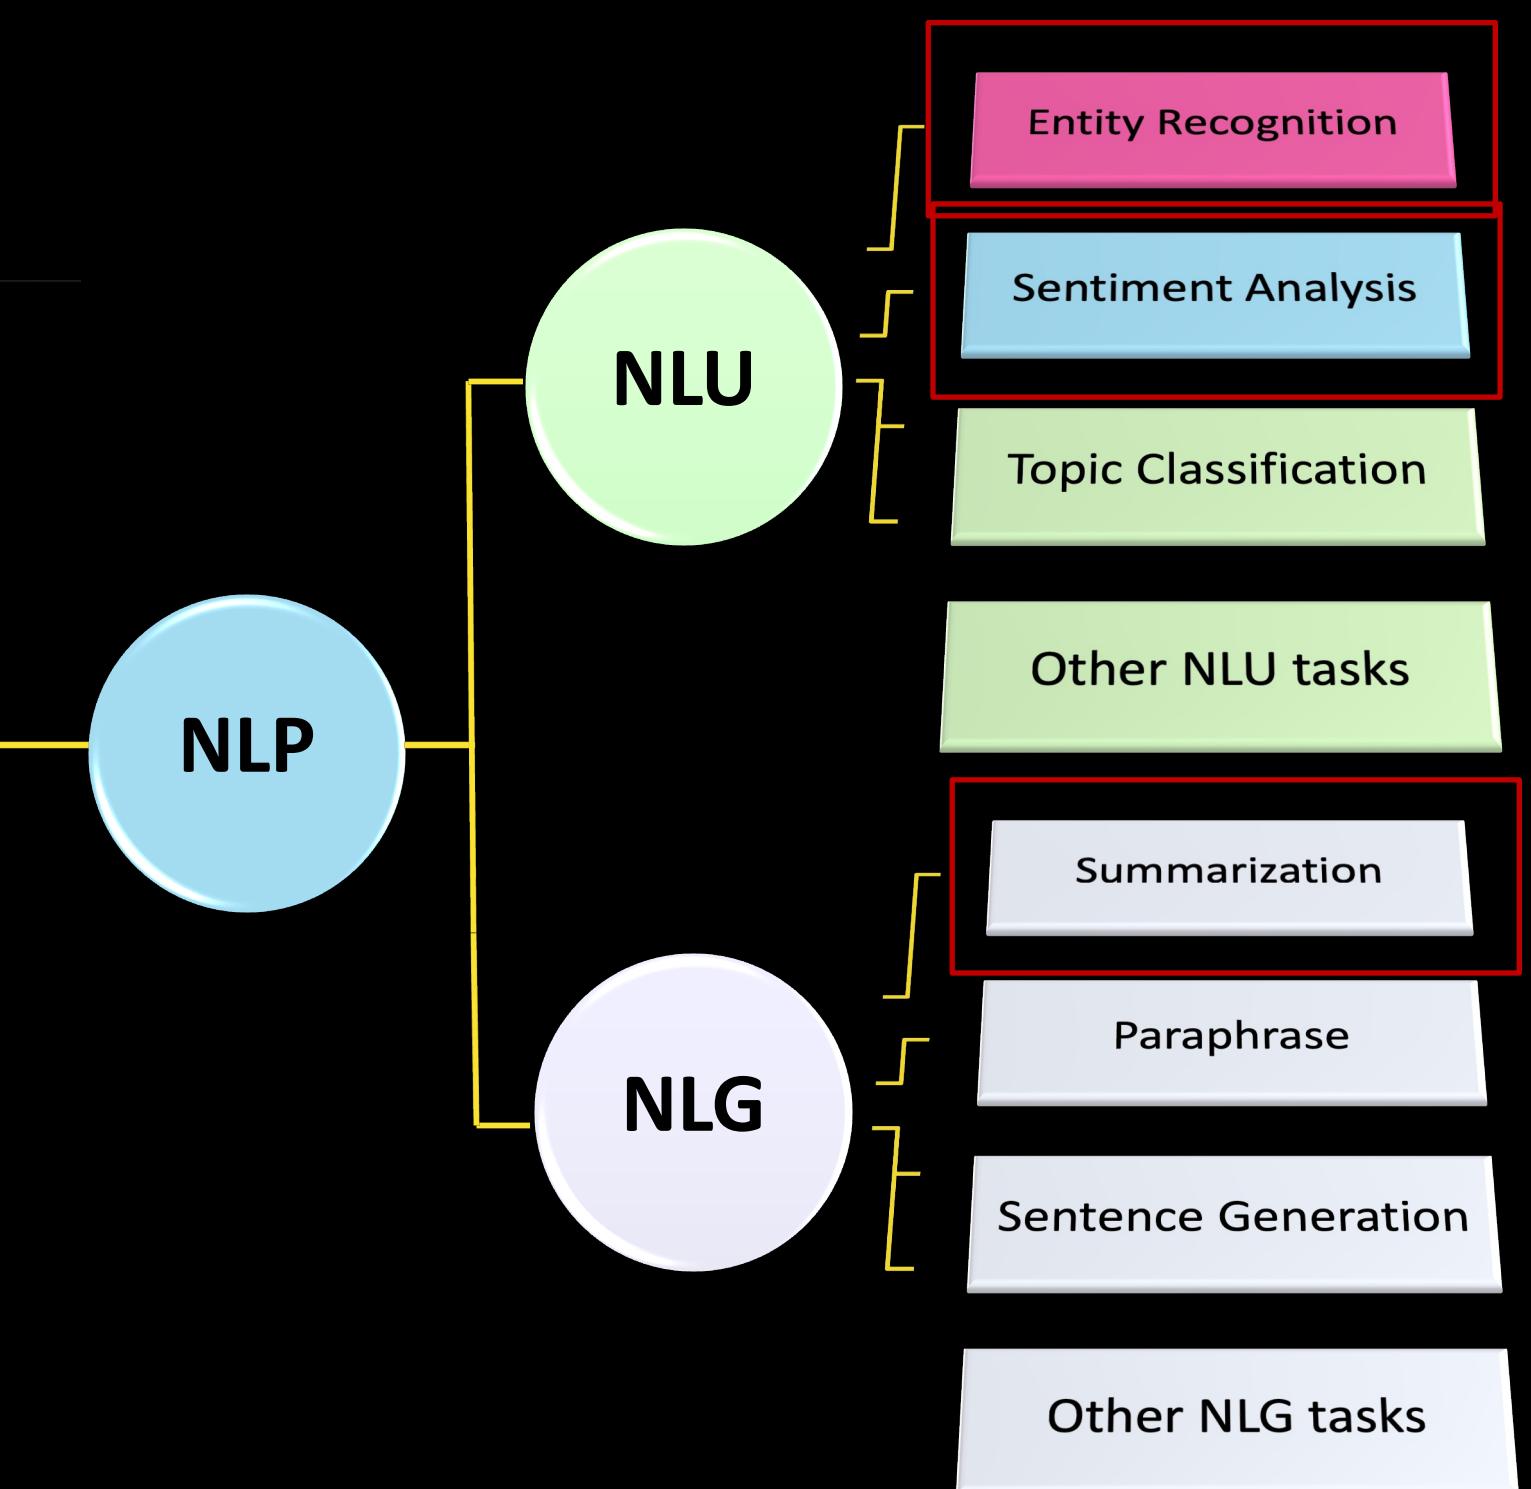


Foundation Model
Large Language Model
GPT-3

Prompt Instruction

Extract the name of this person in this text.

Text: "My name is Simon,
order status?"



Prompt Instruction

Decide whether a phrase's sentiment is positive, neutral, or negative.

Phrase: "How can I help you today?"

Sentiment: Positive

Completion

Entity (Name):

Simon

Completion

Sentiment:

Positive

Summary:

Completion

Customer calling regarding an order.

Prompt Instruction

Summarize the following conversation:

Agent: How can I help you today?

Customer: My name is Simon, order status?

Conversational AI Application

Agent:

How can I help you today?

Customer:

My name is Simon, order status?

Summary of Conversation

Customer calling regarding an order.

Abstractive
Summarization

Zero-Shot
The model predicts the answer given only a natural language description of the task.

One-Shot
In addition to the task description, the model sees a single example of the task

Few-Shot
In addition to the task description, the model sees a few examples of the task.

Text Classification

Named Entity Recognition (Zero-Shot)

Extract the **name** and **mailing address** from this email:

Dear Kaylee,

It was great to talk to you at the seminar. I thought Jane's talk
was quite good.

Thank you for the book. Here's my address 2111 Ash Lane,
Boulder, CO 80301

Best,

Marilú

Name: Marilú

Mailing Address: 2111 Ash Lane, Boulder, CO 80301

Reason Over Unstructured Text

Question and Answering

On Monday, Microsoft launched an OpenAI service as part of its Azure cloud platform, offering businesses and start-ups the ability to incorporate models like ChatGPT into their own systems. The company has already been building AI tools into many of its consumer products, such as a DALL-E 2 feature in its Bing search engine that can create images based on a text prompt, and the information reported recently that it's working to bring more of them to Microsoft Office as well.

What will Microsoft incorporate in its own systems?

Microsoft plans to incorporate AI tools into its own systems, such as a DALL-E 2 feature in its Bing search engine that can create images based on a text prompt, and the information reported recently that it's working to bring more of them to Microsoft Office as well

Summarization

Extractive vs Abstractive

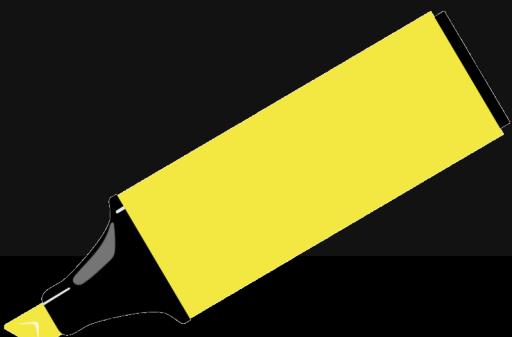
Extractive

Select parts of the original text to form a summary

- Easier
- Restrictive (no paraphrasing)

Peter and Elizabeth took a taxi to attend the night party in the city. While in the party, Elizabeth collapsed and was rushed to the hospital.

Peter and Elizabeth attend party city. Elizabeth rushed hospital



Abstractive (Semantic)

Generates new text using natural language generation techniques

- More difficult
- More flexible (more human)

Peter and Elizabeth took a taxi to attend the night party in the city. While in the party, Elizabeth collapsed and was rushed to the hospital.

Elizabeth was hospitalized after attending a party with Peter.



Abstractive Summarization

Contact Center Summarization

Summarize this for a call center agent:

Agent: Thank you for calling ADM. Who am I speaking to?

Customer: Hello, my name is Peter Smith. I own a small business and have some questions regarding payroll processing.

Agent: Good morning, Peter, before we get started may I ask you a few questions so that I better answer your questions today?

Customer: Thank you that is quite helpful. Are there specific regulations that I need to follow?

Agent: Certain aspects of payroll processing are regulated by the Internal Revenue Service (IRS) and the Department of Labor (DOL)

The agent is speaking to a customer named Peter Smith. Peter is asking for tips on manual payroll processing. The agent also informs Peter that certain aspects of payroll processing are regulated by the Internal Revenue Service (IRS) and the Department of Labor (DOL).

Prompt Insert

Only DaVinci-002+

Before

These are agenda topics for a customer presentation.

1. Introduction to GPT3

[insert]

10. Conclusion

After

These are agenda topics for a customer presentation.

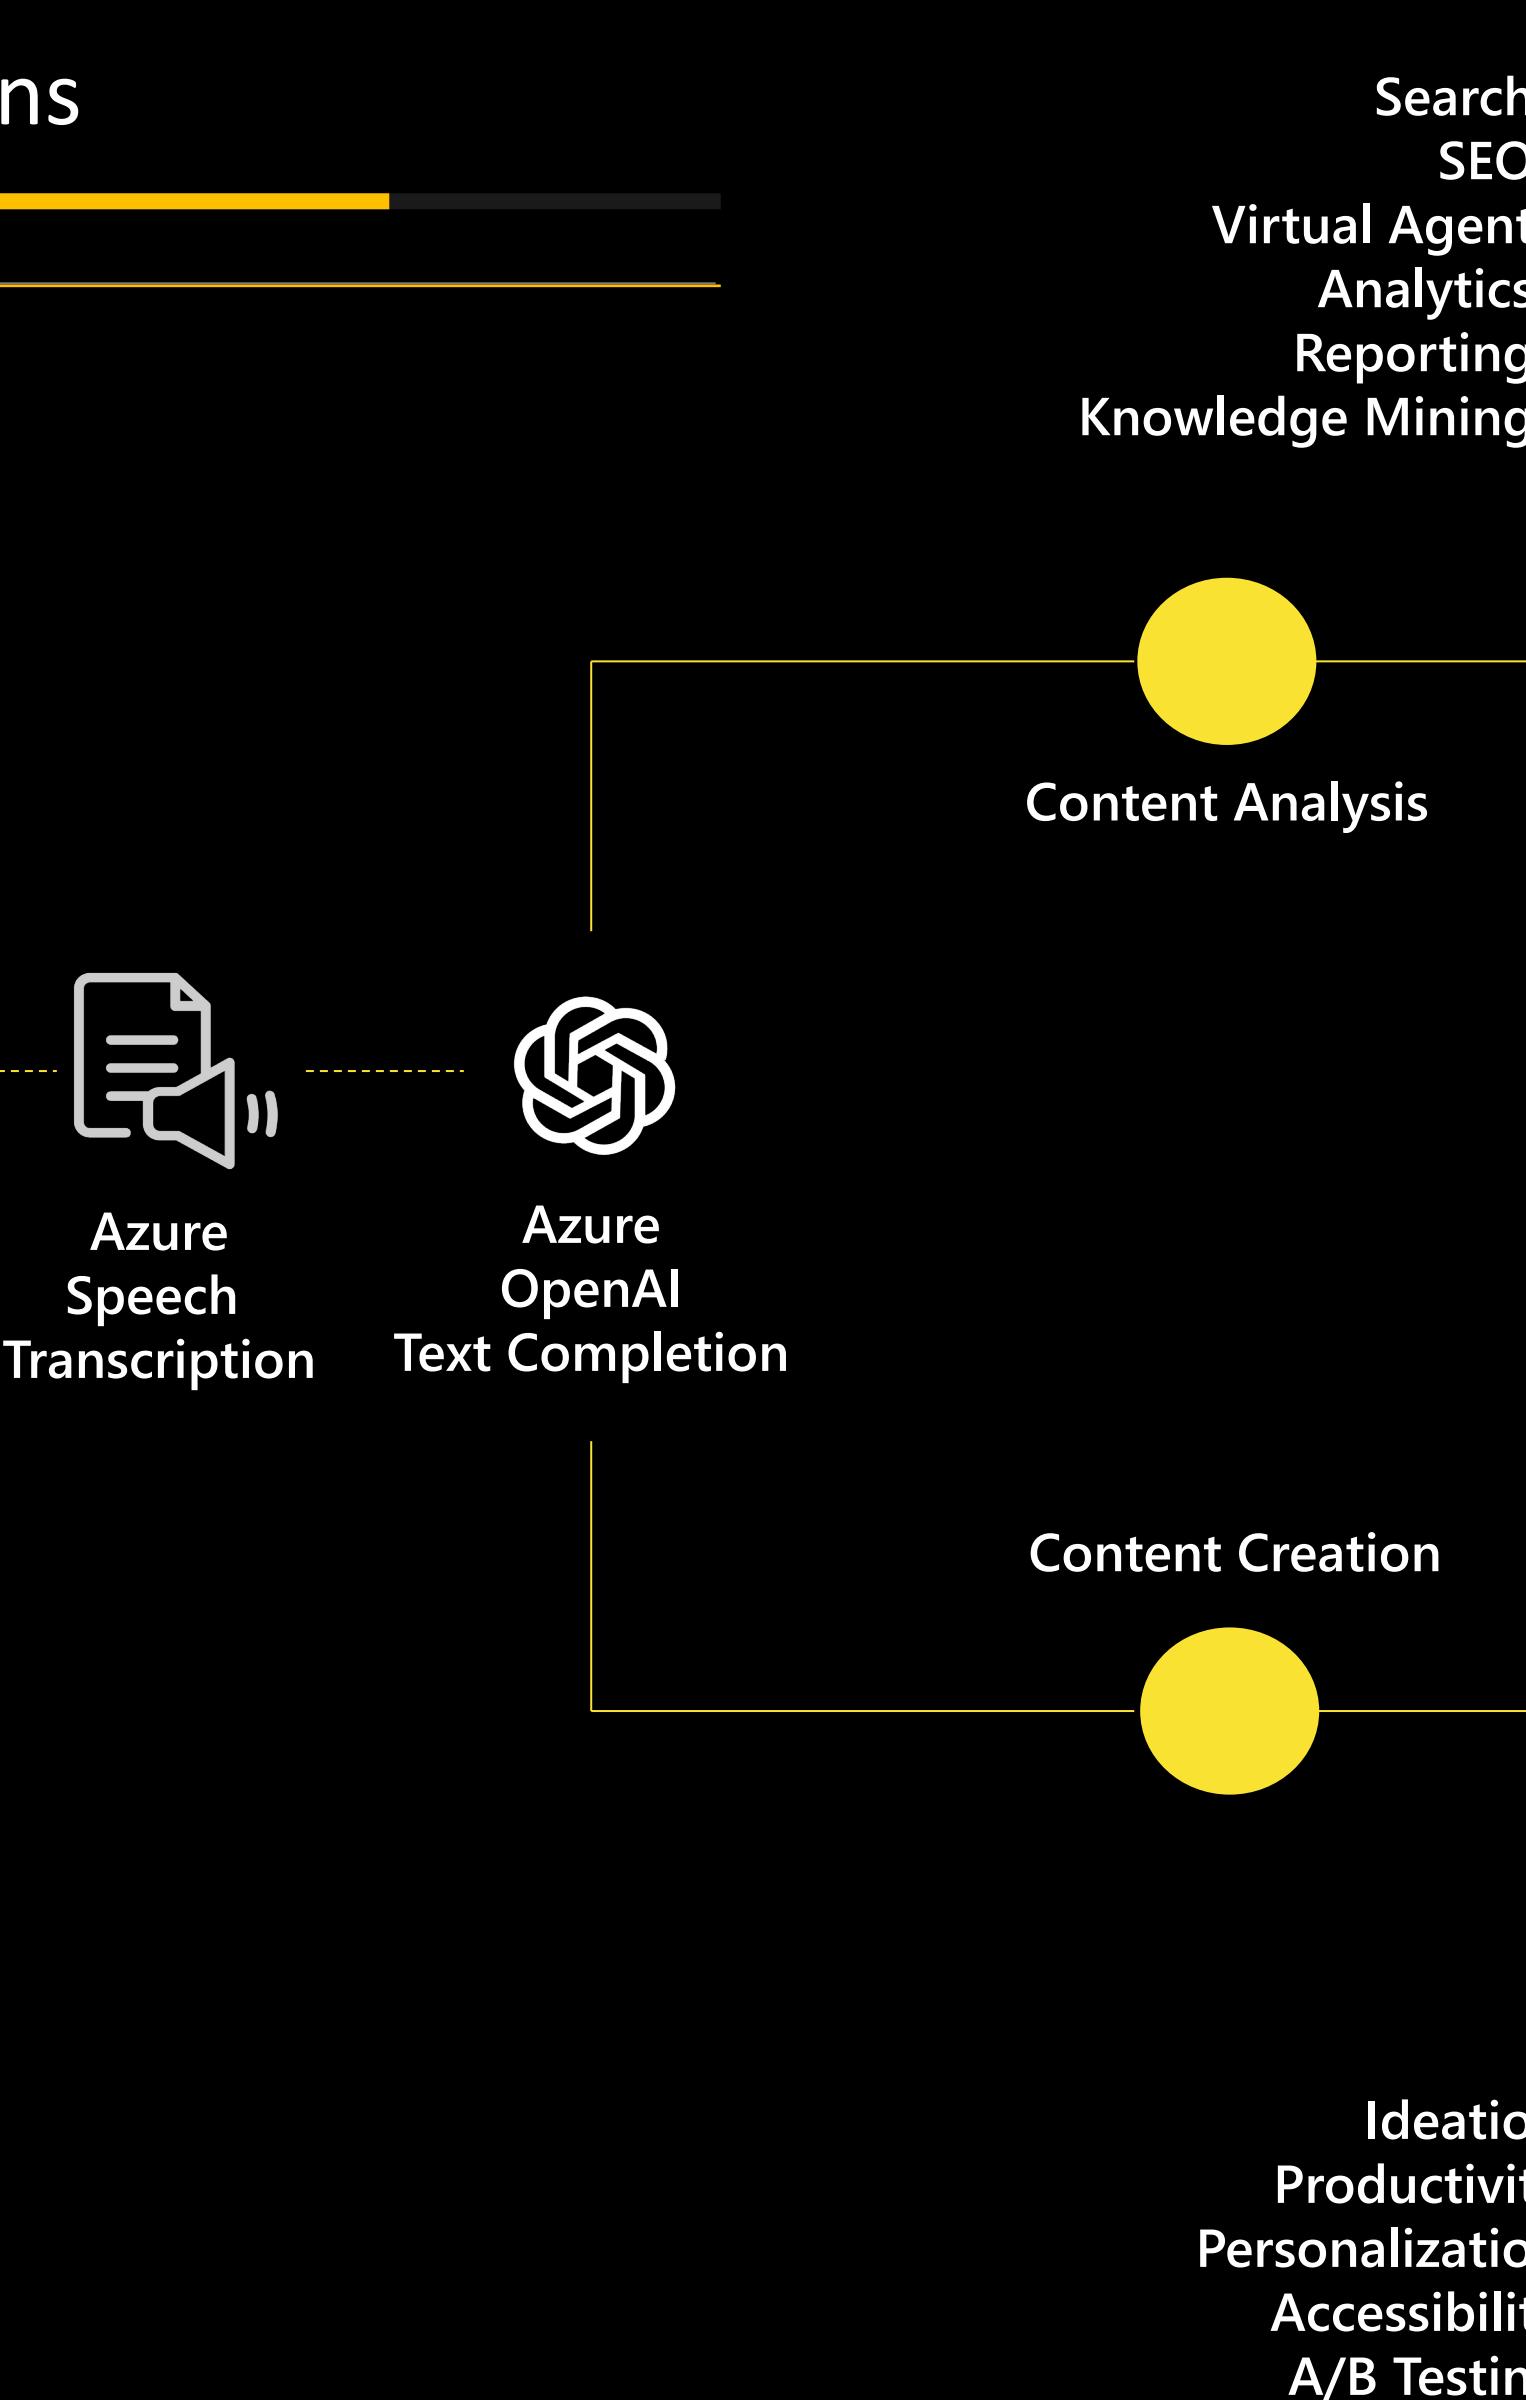
1. Introduction to GPT3
2. Overview of GPT3 Features
3. Benefits of GPT3
4. GPT3 Use Cases
5. GPT3 Pricing
6. GPT3 Security
7. GPT3 Support
8. GPT3 Integration
9. Q&A Session
10. Conclusion

Media Example

News Analyses & Article Creations



News Broadcast (Global Warming)



Topic Classification

Global warming, Deforestation, Carbon footprint

Entity Extraction

Organizations: IPCC, UNFCCC, Green Peace
Geography: Canada, USA

Key Word Extraction

Human activities, fossil fuels, earth atmosphere

Question and Answer

What is the Intergovernmental Panel on Climate Change (IPCC)?

The IPCC is an international organization that studies climate change and the effectiveness ...

Video summarization

The article discusses about global warming and its effects on the Earth's atmosphere, wildlife, and human communities. It states that the primary cause of global warming is

News article generation (or blogs, social media)

Global warming is the gradual increase in the overall temperature of the Earth's atmosphere, primarily caused ...

Script Generation

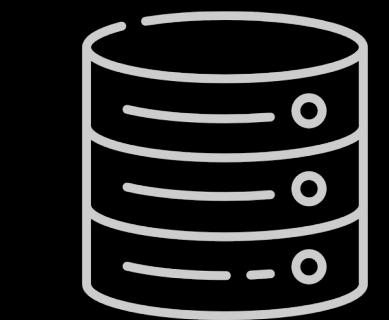
Act 1: The show opens with a shot of a beautiful coastal town
Act 2: As the town struggles to cope with the crisis ...

Personalized Content generation (or Advertising)

Simon, as someone passionate about global warming, you are aware of the urgent threat it poses to our planet ...

Sports & Entertainment Example

Natural language to SQL to surface stats data (Fan Engagement)



Sports Stats
Database



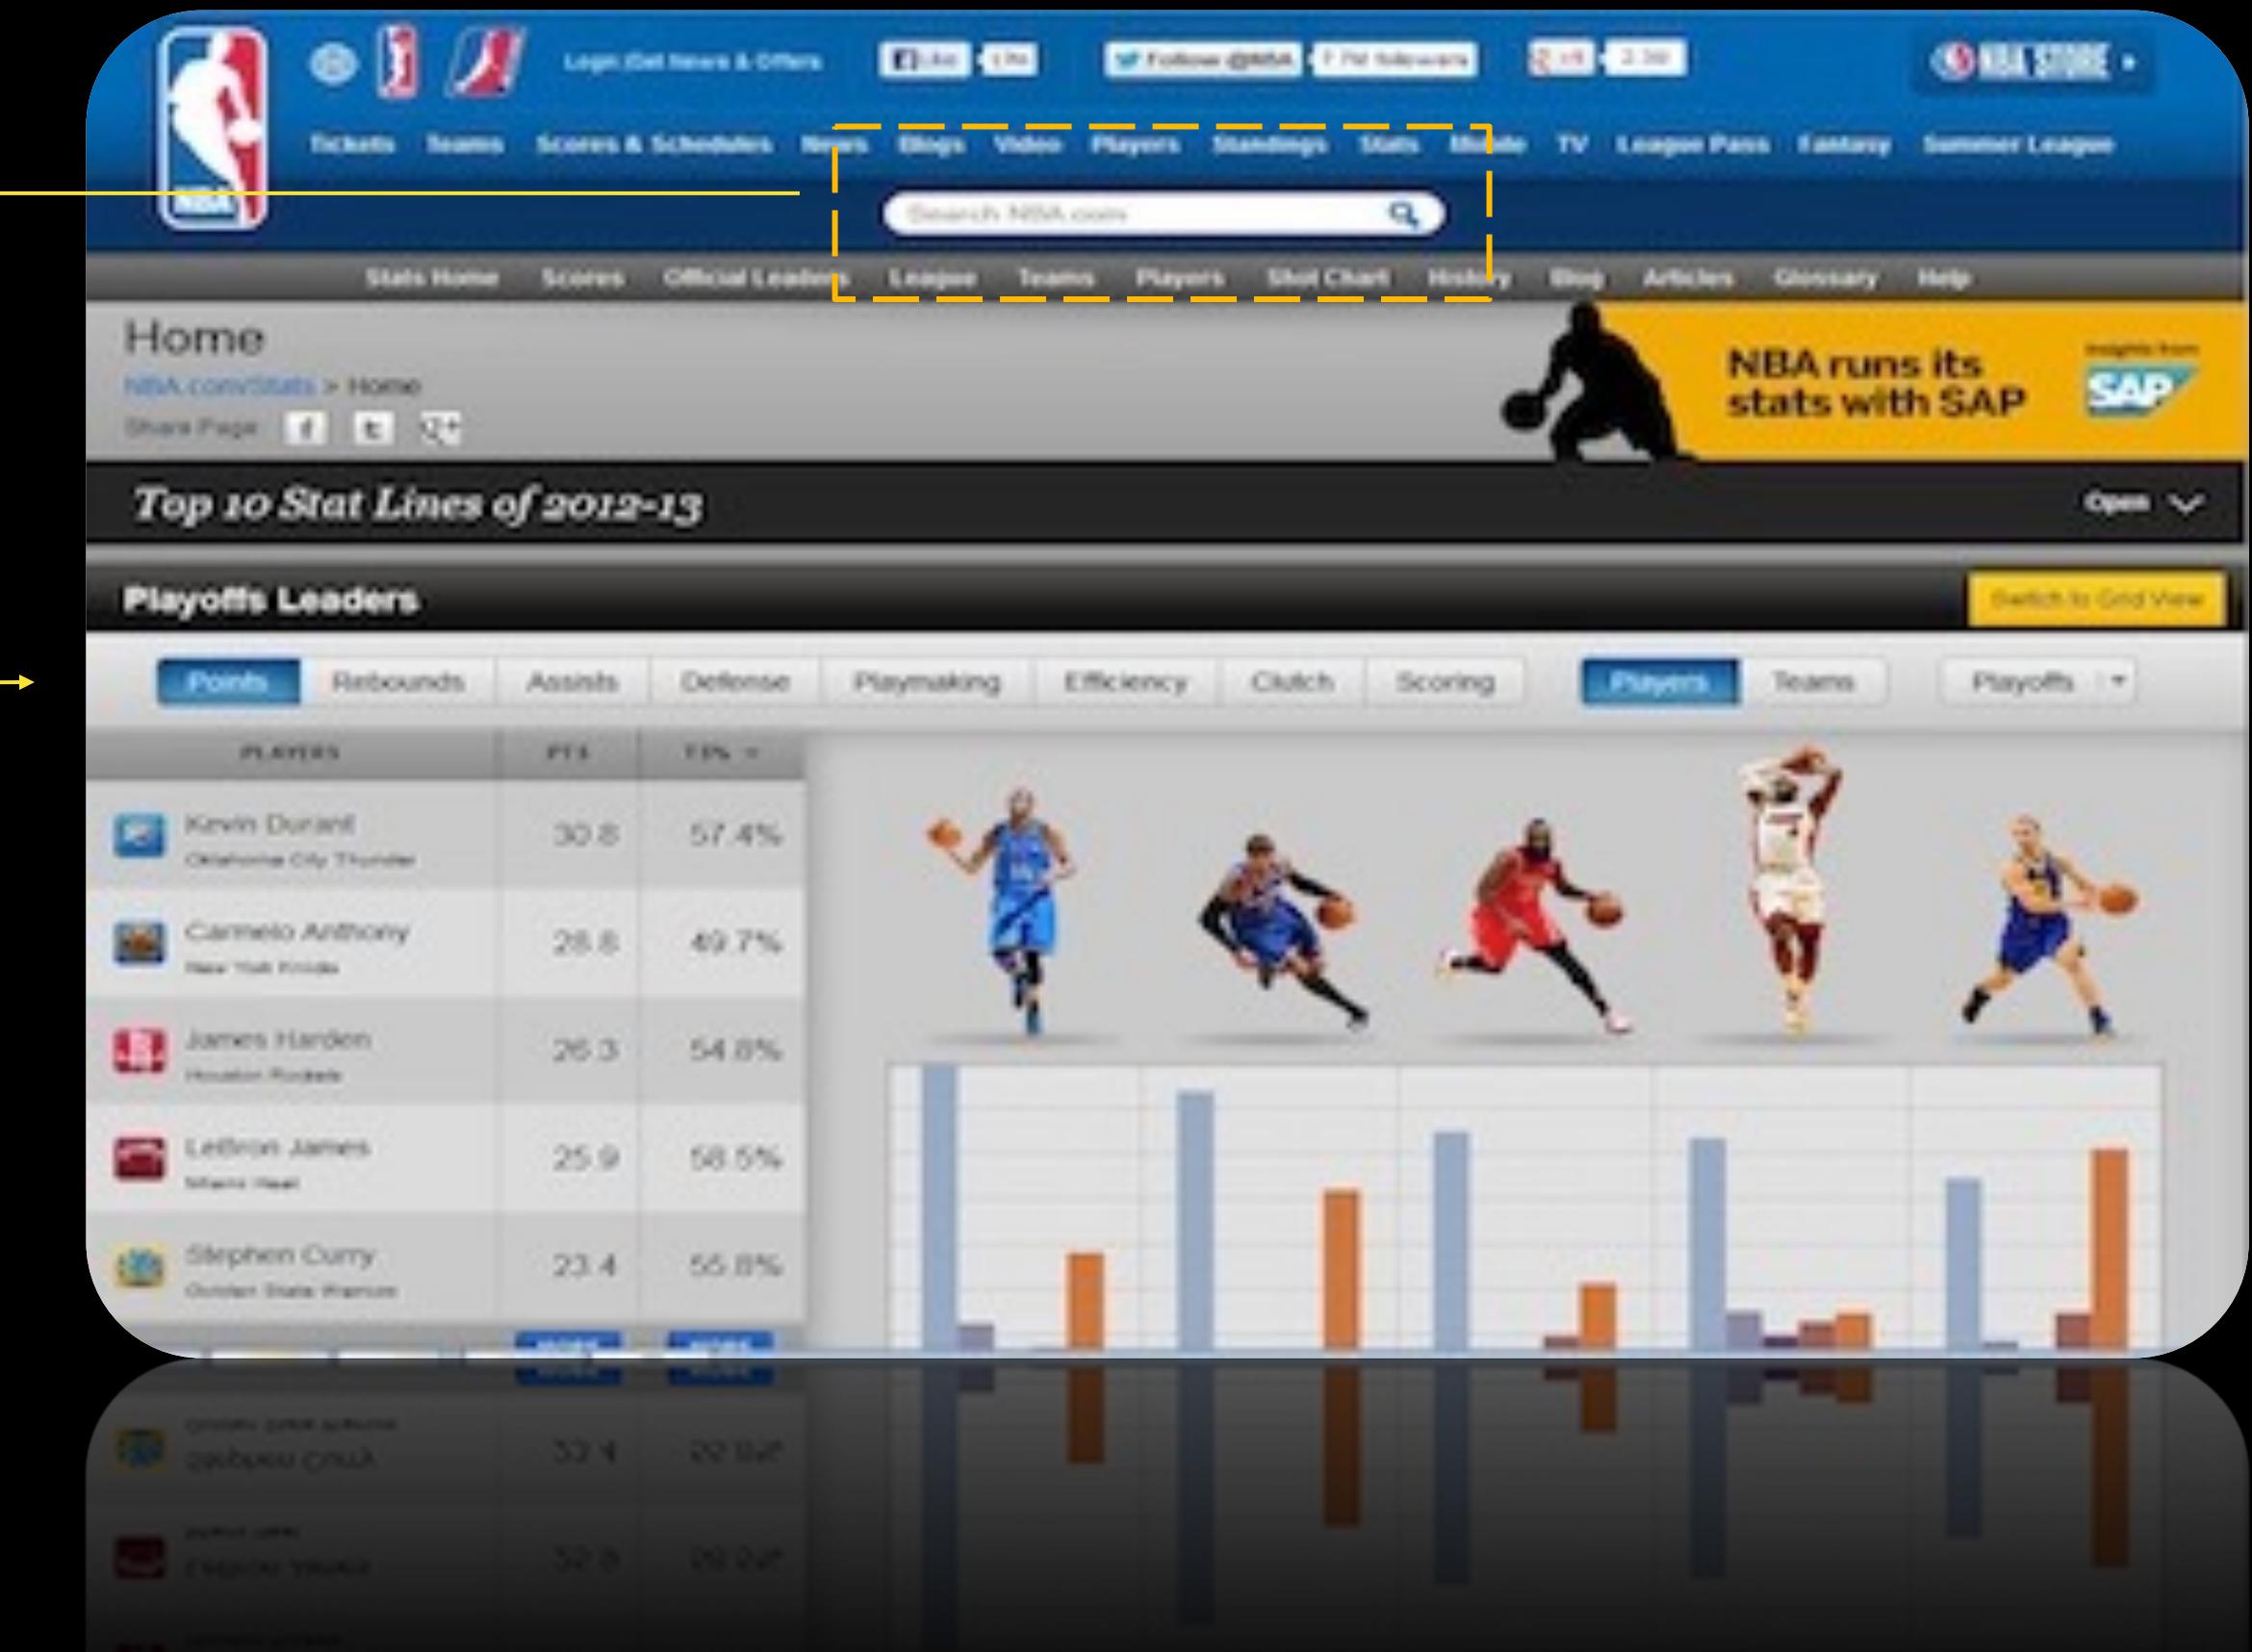
Azure
OpenAI
Codex

"How many points did Lebron James score in 2021?"

```
SELECT SUM(points) FROM basketball_stats WHERE player = 'LeBron James' AND game_date BETWEEN '2021-01-01' AND '2021-12-31';
```

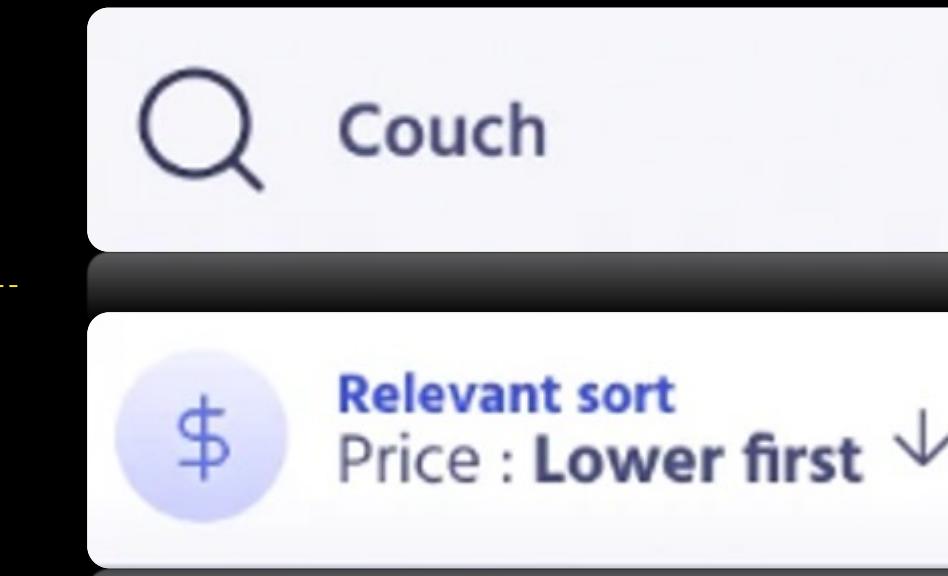
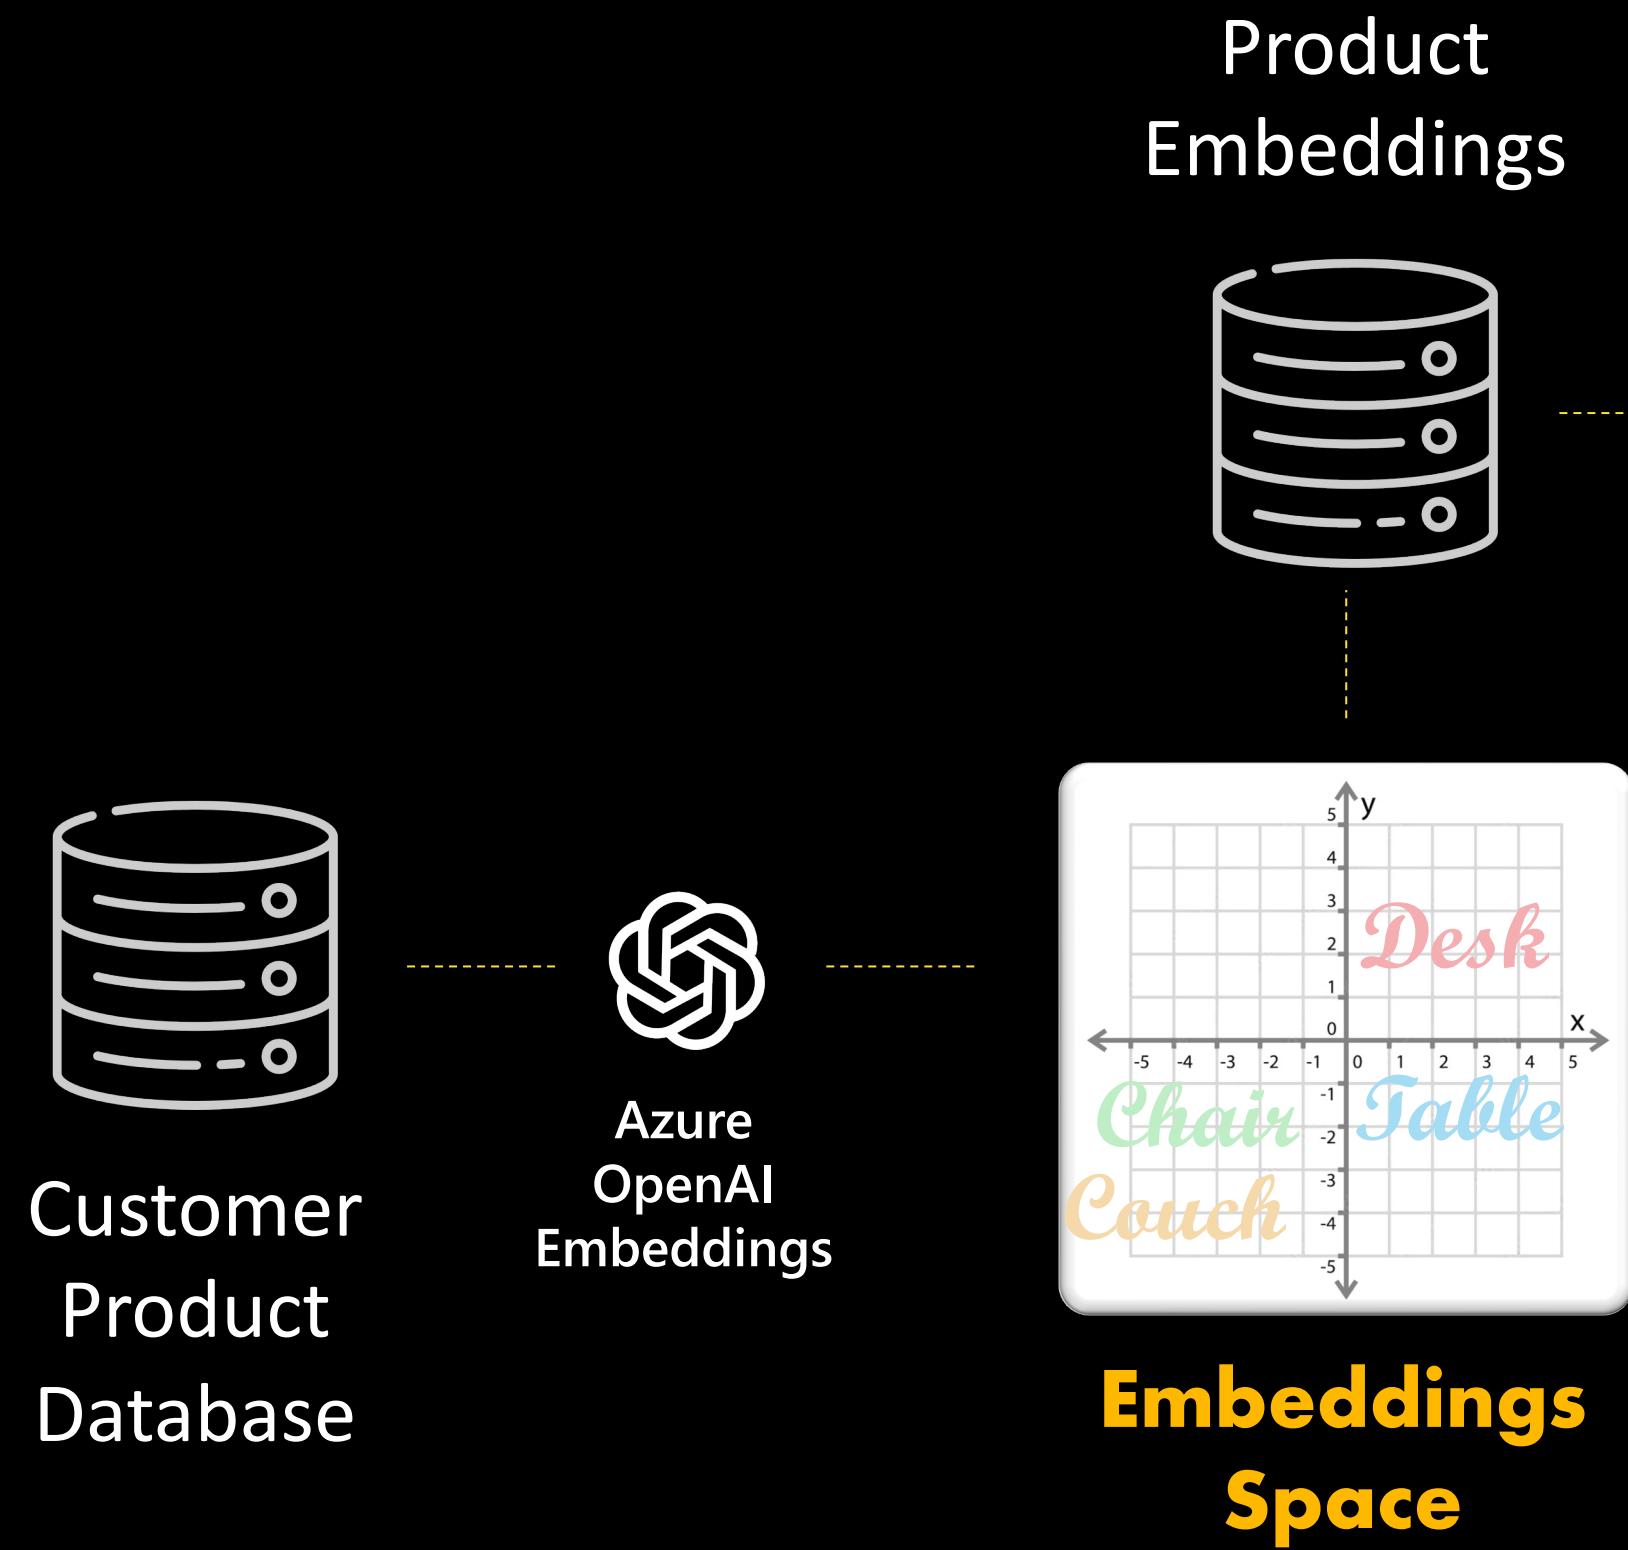
basketball_stats (player, team, points, rebounds, assists, steals, blocks, turnovers, games_played)
("LeBron James", "Los Angeles Lakers", 25, 7, 10, 1, 0, 3, 82), ("Kevin Durant", "Brooklyn Nets", 30, 7, 4, 1, 2, 3, 72)

basketball_games (home_team, away_team, home_score, away_score, game_date)
("Los Angeles Lakers", "Brooklyn Nets", 110, 105, '2021-12-01'),
("Milwaukee Bucks", "Golden State Warriors", 120, 115, '2021-12-02')

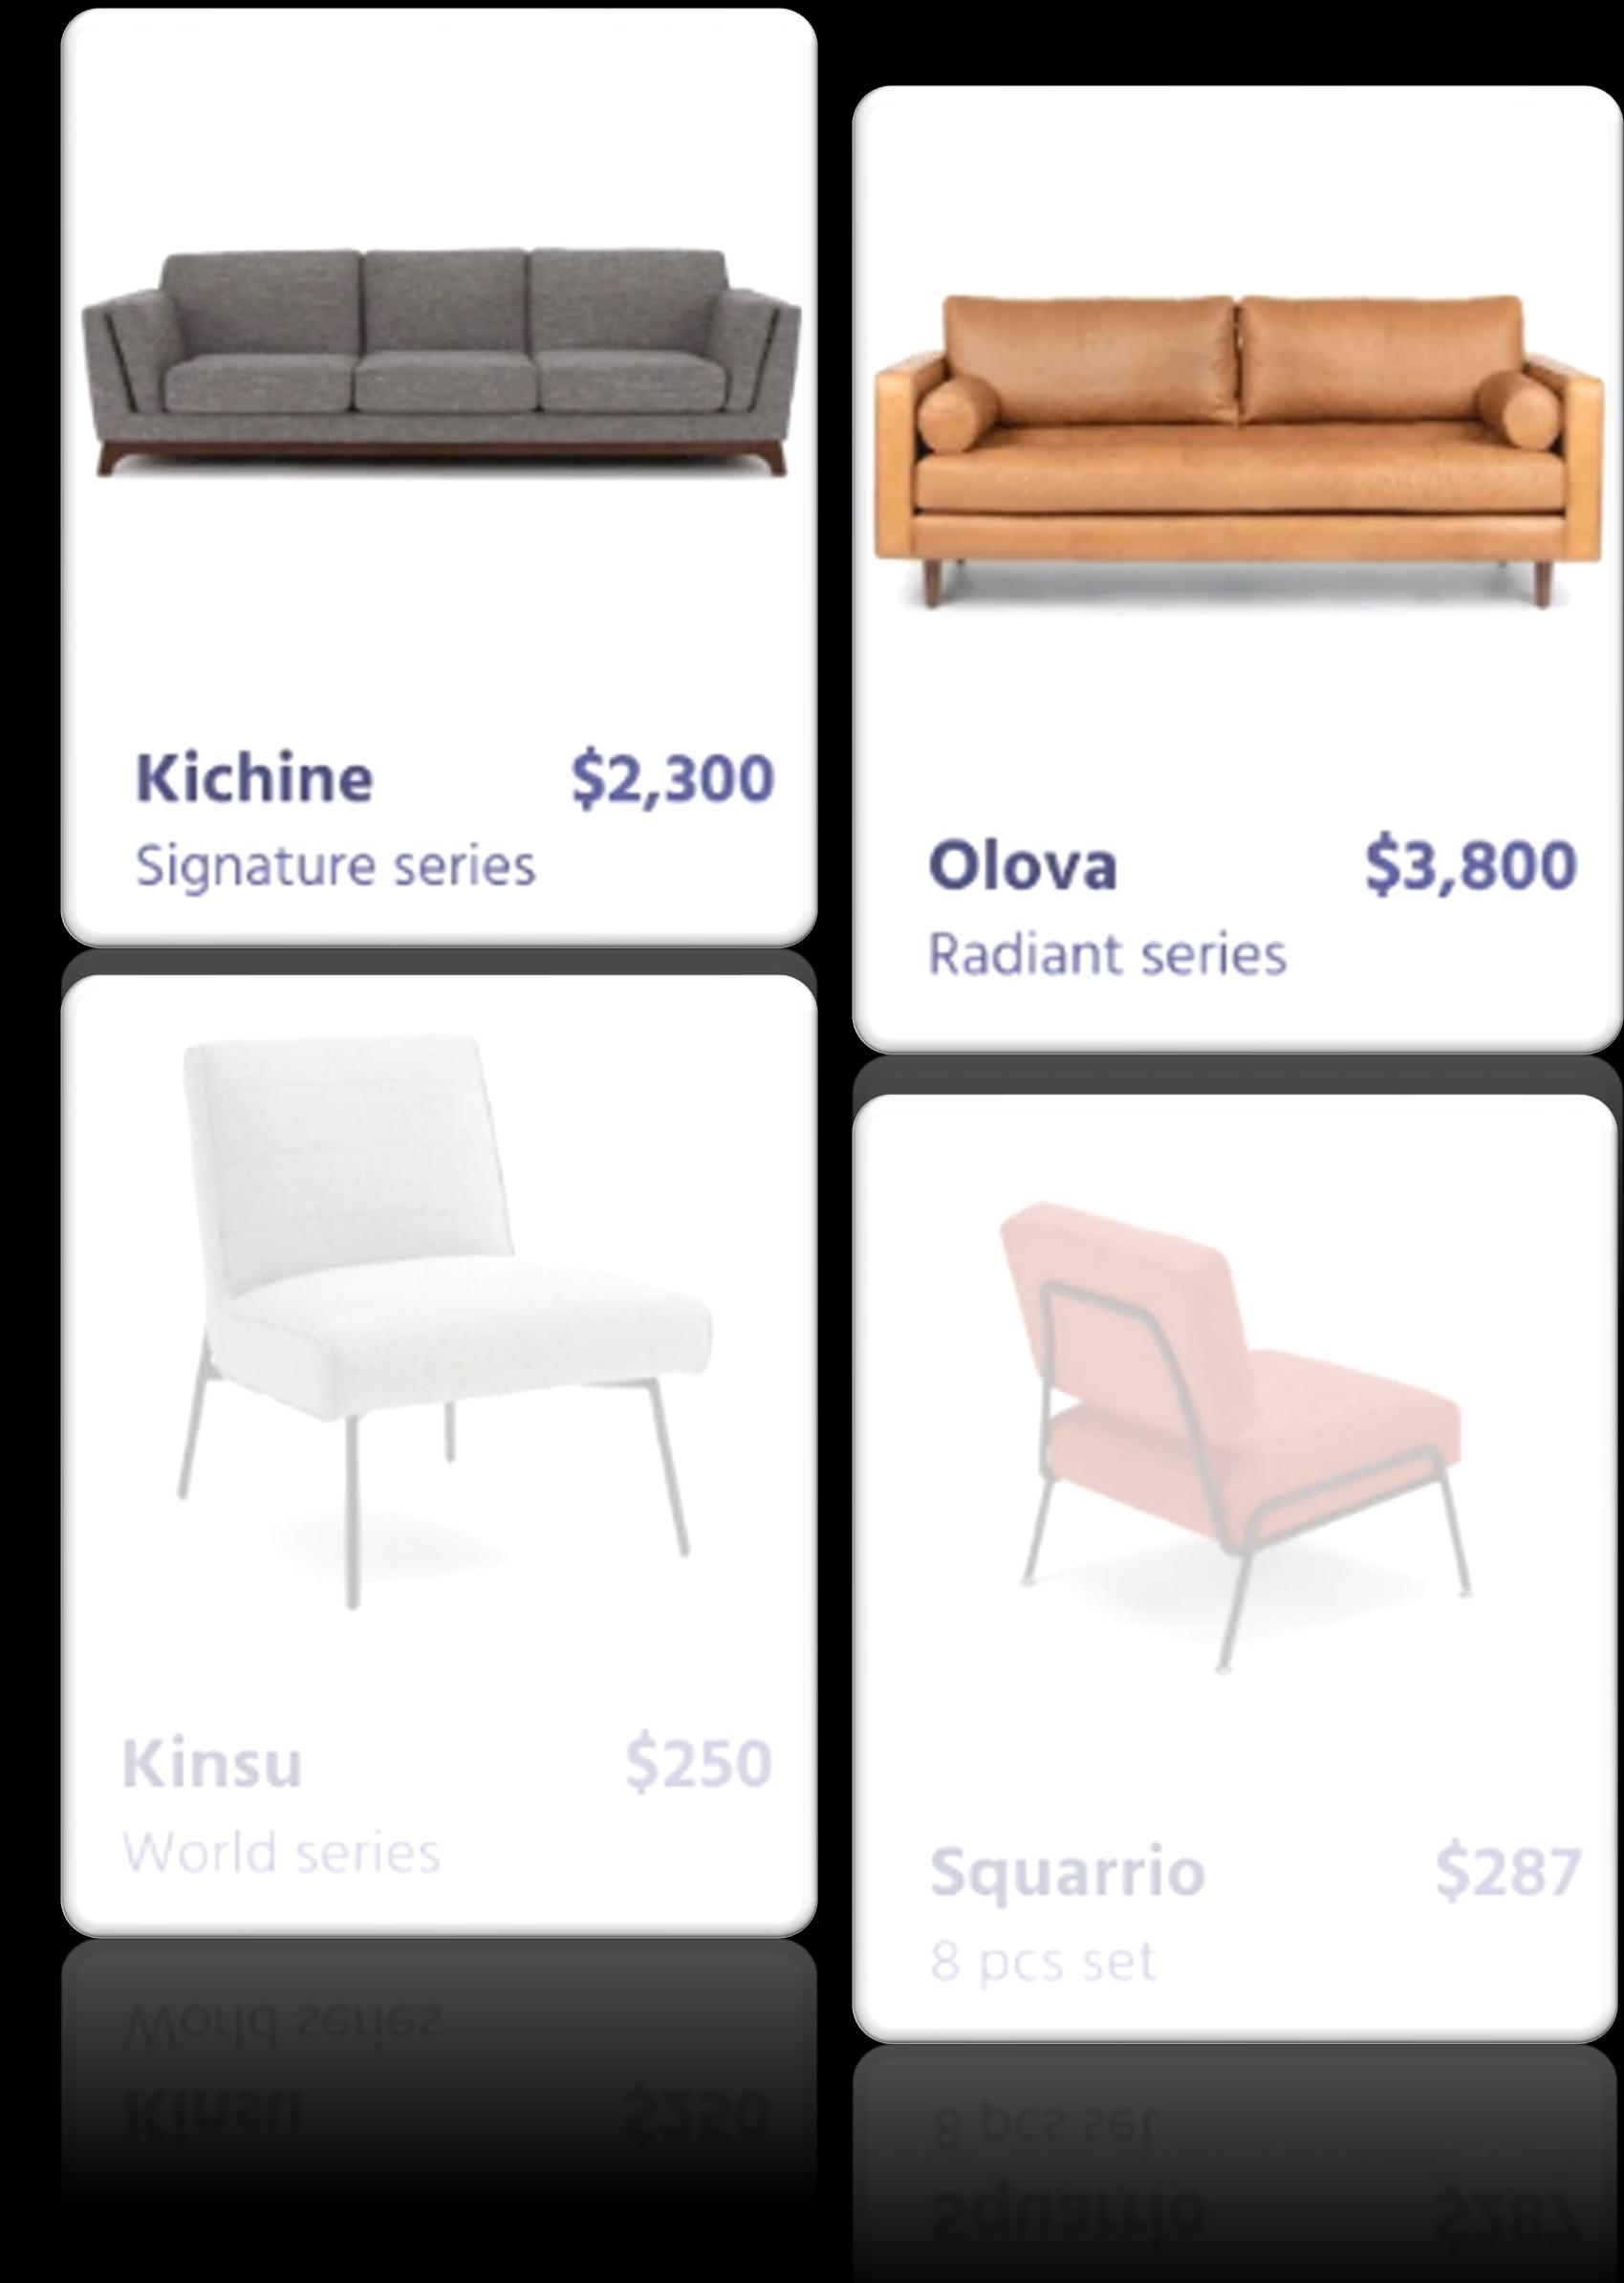


Retail Example

Similarity Search



Similarity Search



Marketing Example

Digital Creative Assistant – Dalle-2

Synthetic Brand Ambassador (GDPR-safe)

Input
Generate a white female

Generate

Output

No code Web and App Development

Input

Output

```
constructor(props) {  
  super(props);
```

Add \$3 Withdraw \$5
Give away all my money

My balance is 9

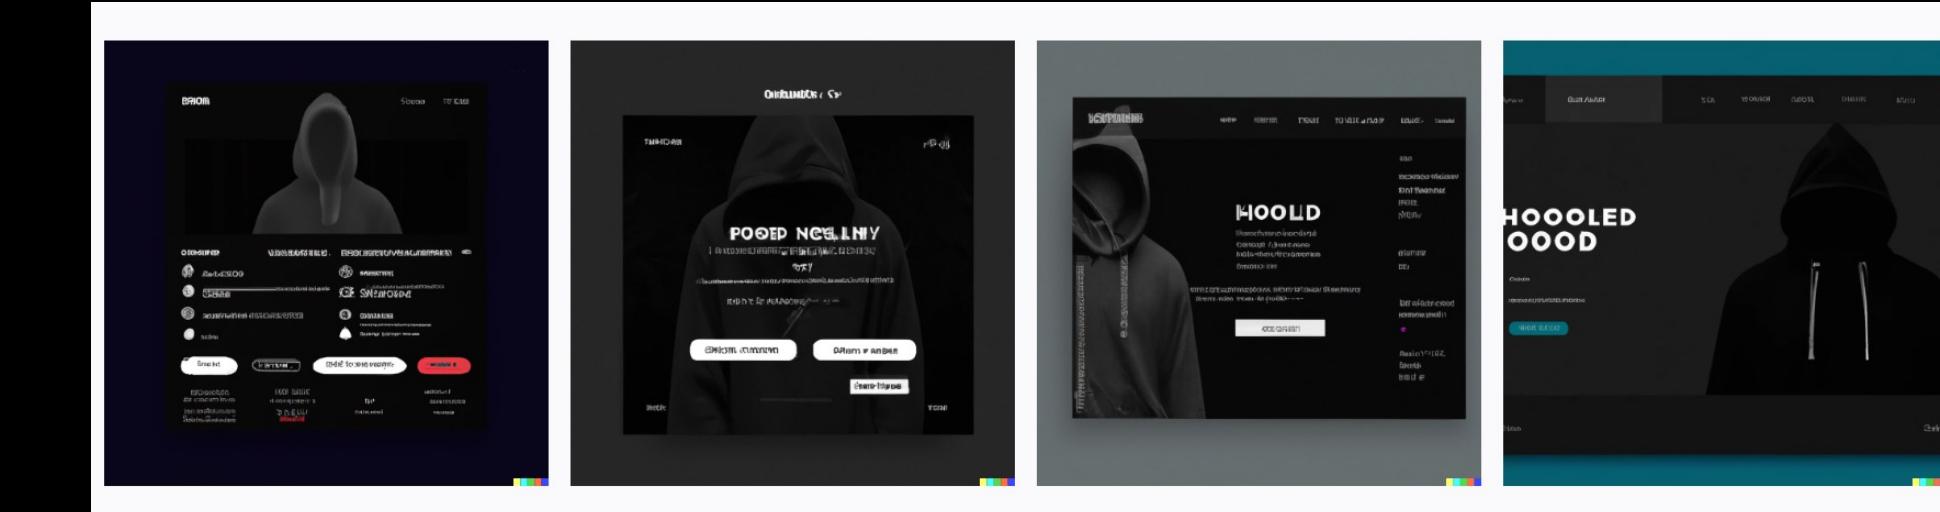
Generate images, concepts and ideas

Input

a green c

Output

Produce rough Layouts



- Intros
- ChatGPT 101
- Business Examples
- Exercise

AWS Bedrock Demo



Text generation



Chatbot



Summarization



Question
answering



Image
generation



Code
generation

AWS Bedrock Demo

Amazon Bedrock > Text playground

Text playground Info

Streaming

AI Anthropic ▾ Claude V2 v2

Using on-demand capacity. [Switch Provisioned Throughput](#)

Please rewrite the above paragraph to make it understandable to a 5th grader.

Please output your rewrite in <rewrite></rewrite> tags.

Assistant:

```
<rewrite>
```

Dear Support Team,

I got some feedback from a customer about an issue they are having when printing shading charts on their printers.

They are getting an error message (error code 3.1.113) whenever they try to print a shading chart using their computer. This is happening on two of their printers - Engine1 and Engine2.

They have tried turning the printers off and back on again and restarting the software, but they still get the same error message when trying to print a shading chart.

They also tried reloading the firmware (the software that runs on the printer) for one part of Engine2, but that didn't fix the problem either.

I attached some technical files and log data to a file sharing website for you to look at and try to understand the issue better.

Please take a look and let me know if you have any ideas on how to fix this printing problem.

Thank you!

Sincerely,

Inference configuration Reset

▼ Randomness and diversity Info

Temperature: 0.5

Top P: 1

Top K: 250

▼ Length Info

Maximum length: 2048

Stop sequences: Add Human:

The screenshot shows the AWS Bedrock Text playground interface. At the top, there are dropdown menus for 'Anthropic' and 'Claude V2 v2'. A checkbox for 'Streaming' is checked. Below this, there's a large text area containing a message from a customer about printer errors. To the right of the message is an 'Inference configuration' sidebar with sliders for Temperature (0.5), Top P (1), and Top K (250), and a text input for Maximum length (2048). The sidebar also includes sections for 'Randomness and diversity' and 'Length', and a 'Stop sequences' section with a text input and an 'Add' button. The main text area has a blue border and contains several paragraphs of text, some of which are redacted with '<rewrite>' tags.

AWS Bedrock Demo

Titan Text G1 - Express v1 - preview

Action items from a meeting transcript

This prompt creates a list of action items from a meeting transcript.

Summarization

Jurassic-2 Ultra v1

Article Summarization

An example prompt that summarizes a news article in bullet points

Summarization

Claude v2

Content Generation

An example prompt to generate a paragraph from instructions

Text generation

Command v14.7

Contract Entity Extraction

Use generative models to easily parse out entities from text

Open ended text generation

Stable Diffusion XL v0.8 - preview

Create an image

This prompt creates an image based on a short description

Image gen

Titan Text G1 - Express v1 - preview

Creating a table of product descriptions

This is a prompt for generating product descriptions that incorporate keywords

Open ended text generation

Jurassic-2 Mid v1

Extract topics and sentiments from reviews

An example prompt that extracts topics from hotel reviews and determines the sentiment

Information extraction

Jurassic-2 Mid v1

Generate product descriptions

An example prompt that generates a product description for an eCommerce website

Text generation

Claude v2

Information extraction

An example prompt to extract information from passages

Information extraction

Claude v2

Multiple choice classification

An example prompt to classify emails

Question answering

Jurassic-2 Ultra v1

Outline generation

An example prompt that generates outline for a blog post based on a title

Text generation

Command v14.7

Product Description to Benefits

Turn a product description into a list of functional, emotional, and social benefits

Open ended text generation

AWS Bedrock Demo

Claude v2

Question and answer

An example prompt for simple Q&A.

Question answering

Claude v2

Removing PII

An example prompt to remove personally identifiable information (PII) from given text

Information removal

Titan Text G1 - Express v1 - preview

Summarize the key takeaways

This prompt summarizes a document into the main takeaways

Summarization

Command v14.7

Transcript Summarization

Extract the main point of a conversation

Open ended text generation

Titan Text G1 - Express v1 - preview

Write an article

This is a prompt that generates an article for kids on a technical topic

Open ended text generation

Titan Text G1 - Express v1 - preview

Writing a promo doc

This prompt takes simple bulletpoints of accomplishments and creates a promotion narrative for an employee

Rewrite

AWS Bedrock Demo

Generation

- [Simple use case with boto3](#): In this notebook, you generate Amazon Titan model directly with boto3
- [Simple use case with LangChain](#): We then perform the same task using LangChain
- [Generation with additional context](#): We then take this further by adding context to the prompt to improve the response.

Summarization

- [Small text summarization](#): In this notebook, you use use Bedrock API to generate summaries for small amount of text.
- [Long text summarization](#): The above approach may not work well for long texts due to token limit. This notebook shows an approach of breaking the text into smaller chunks, summarizing each chunk, and then summarizing the summaries.

Question Answering

- [Simple questions with context](#): This notebook shows a simple question answering system using a Language Model directly.
- [Answering questions with Retrieval Augmented Generation](#): This notebook shows an architecture called Retrieval Augmented Generation (RAG). It uses a pre-trained Language Model (parametric) and augments the prompts by adding the relevant context from a database.

Chatbot

- [Chatbot using Claude](#): This notebook shows a chatbot using Claude API.
- [Chatbot using Titan](#): This notebook shows a chatbot using Amazon Titan model.

Text to Image

- [Image Generation with Stable Diffusion](#): This notebook demonstrates image generation with using the Stable Diffusion model

Code Generation, SQL Generation, Code Translation and Explanation

1. [Code Generation](#): Demonstrates how to generate Python code using Natural language. It shows examples of prompting to generate simple functions, classes, and full programs in Python for Data Analyst to perform sales analysis on a given Sales CSV dataset.
2. [Database or SQL Query Generation](#) : Focuses on generating SQL queries with Amazon Bedrock APIs. It includes examples of generating both simple and complex SQL statements for a given data set and database schema.
3. [Code Explanation](#) : Uses Bedrock's foundation models to generate explanations for complex C++ code snippets. It shows how to carefully craft prompts to get the model to generate comments and documentation that explain the functionality and logic of complicated C++ code examples. Prompts can be easily updated for another programming languages.
4. [Code Translation](#) : Guides you through translating C++ code to Java using Amazon Bedrock and LangChain APIs. It shows techniques for prompting the model to port C++ code over to Java, handling differences in syntax, language constructs, and conventions between the languages.

Entity Extraction

- [Entity Extraction with Claude v2](#): This notebook shows how LLM can be used to extract specific information from natural text.

Chatbot and LLMs Guardrails

- [LLM & NeMo Guardrails](#): Explores the implementation of guardrails for Language Model (LLM) generated responses using Amazon Bedrock and NVIDIA's NeMo. It highlights the utility of guardrails in ensuring responses adhere to desired parameters, providing a more advanced mechanism over standard system prompts. This notebook demonstrates the integration and configuration of guardrails with NeMo and Bedrock, showcasing various guardrail configurations like Jailbreaking Rail, Topical Rail, Moderation Rail and Fact Checking for safer and more reliable AI interactions.

AWS Bedrock Models

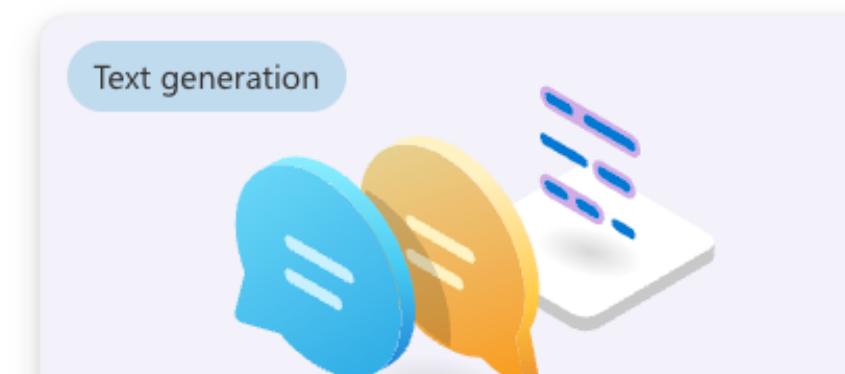
Base models (9)		
Models	Access status	Modality
AI21 Labs	In Progress	Text
Jurassic-2 Ultra	In Progress	Text
Amazon	In Progress	Text
Titan Embeddings G1 - Text	Access granted	Embedding
Anthropic	Unavailable	Text
Titan Text G1 - Lite Preview	Unavailable	Text
Cohere	Unavailable	Text
Claude	Use case details required	Text
Stability AI	Use case details required	Text
Command	In Progress	Text
Stable Diffusion XL Preview	In Progress	Image

Class exercise MSFT OpenAI Studio

Welcome to Azure OpenAI service

Explore the generative AI models, craft unique prompts for your use cases, and fine-tune select models.

Get started



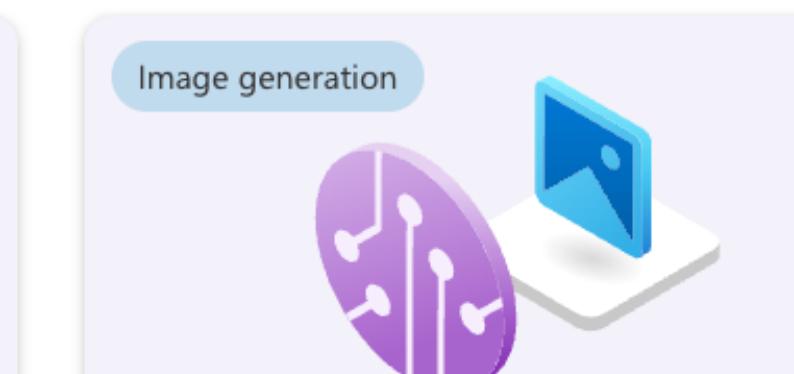
Chat playground

Design a customized AI assistant using ChatGPT. Experiment with GPT-3.5-Turbo and GPT-4 models.

[Try it now](#)

Completions playground

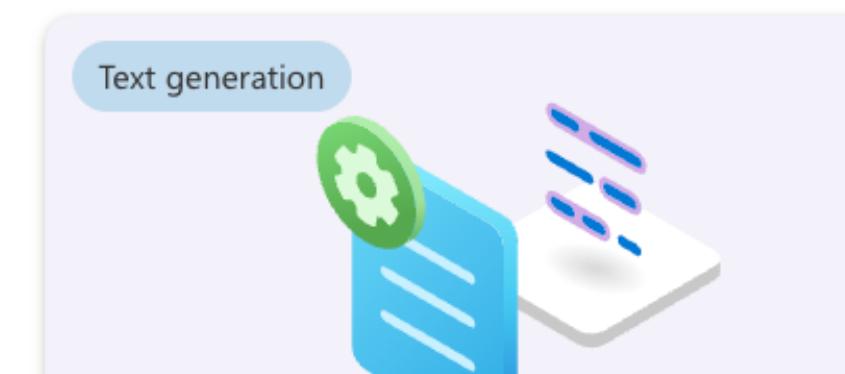
Experiment with completions models for use cases such as summarization, content generation, and classification.

[Try it now](#)

DALL-E playground

[PREVIEW](#)

Generate unique images by writing descriptions in natural language.

[Try it now](#)

Bring your own data

[PREVIEW](#)

Connect and ground your data. Deploy to a web app or Power Virtual Agent bot.

[Try it now](#)

Try out common examples

Customer support agent

Get started with a sample chat setup for a customer support agent scenario

[Try it now](#)

Writing assistant

Get started with a sample chat setup for marketing content creation

[Try it now](#)

Summarize an article

Get started with abstractive summarization of a news article

[Try it now](#)

Create cover art

Unleash your creativity by generating creative images for cover art and other commercial use cases

[Try it now](#)

Class exercise

MSFT OpenAI Studio

Azure AI | Azure OpenAI Studio

Azure AI Studio > Completions playground

Privacy & cookies

Parameters

Temperature 1

Max length (tokens) 100

Stop sequences

Top probabilities 0.5

Frequency penalty 0

Presence penalty 0

Pre-response text

Post-response text

Learn more

Completions playground

Deployments Examples

gpt-35-turbo Load an example

Start typing here

Summarize issue resolution from co...

Summarize key points from financial...

Summarize an article (abstractive)

Generate product name ideas

Generate an email

Generate a product description (bul...

Generate a listicle-style blog

Generate a job description

Generate a quiz

Classify Text

Classify and detect intent

Cluster into undefined categories

Analyze sentiment with aspects

Extract entities from text

Parse unstructured data

Translate text

Natural Language to SQL

Natural language to Python

Explain a SQL query

Question answering

Generate insights

Chain of thought reasoning

Chatbot

View code

Azure OpenAI

Playground Chat Completions DALL-E (Preview)

Management Deployments Models Data files Quotas Content filters (Preview)

Class exercise

MSFT OpenAI Studio

Azure AI | Azure OpenAI Studio

Completions playground

Deployments Examples

gpt-35-turbo Summarize an article (abstractive) View code

Provide a summary of the text below that captures its main idea.

At Microsoft, we have been on a quest to advance AI beyond existing techniques, by taking a more holistic, human-centric approach to learning and understanding. As Chief Technology Officer of Azure AI Cognitive Services, I have been working with a team of amazing scientists and engineers to turn this quest into a reality. In my role, I enjoy a unique perspective in viewing the relationship among three attributes of human cognition: monolingual text (X), audio or visual sensory signals, (Y) and multilingual (Z). At the intersection of all three, there's magic—what we call XYZ-code as illustrated in Figure 1—a joint representation to create more powerful AI that can speak, hear, see, and understand humans better. We believe XYZ-code will enable us to fulfill our long-term vision: cross-domain transfer learning, spanning modalities and languages. The goal is to have pre-trained models that can jointly learn representations to support a broad range of downstream AI tasks, much in the way humans do today. Over the past five years, we have achieved human performance on benchmarks in conversational speech recognition, machine translation, conversational question answering, machine reading comprehension, and image captioning. These five breakthroughs provided us with strong signals toward our more ambitious aspiration to produce a leap in AI capabilities, achieving multi-sensory and multilingual learning that is closer in line with how humans learn and understand. I believe the joint XYZ-code is a foundational component of this aspiration, if grounded with external knowledge sources in the downstream AI tasks. We are excited to share our progress and look forward to working with the community to advance AI beyond existing techniques.

Microsoft is working on a new approach to AI that is more holistic and human-centric. The company's Chief Technology Officer of Azure AI Cognitive Services has been working with a team of scientists and engineers to create a joint representation of monolingual text, audio or visual sensory signals, and multilingual text. This representation, called XYZ-code, will enable the development of pre-trained models that can jointly learn representations to support a broad range of downstream AI tasks. Microsoft has already achieved human performance on benchmarks in conversational speech recognition, machine translation, conversational question answering, machine reading comprehension, and image captioning. The company is now working on achieving multi-sensory and multilingual learning that is closer in line with how humans learn and understand. The joint XYZ-code is a foundational component of this aspiration.

How does Microsoft's approach to AI differ from existing techniques?

Microsoft's approach to AI is more holistic and human-centric than existing techniques. The company is working on a joint representation of monolingual text, audio or visual sensory signals, and multilingual text called XYZ-code. This representation will enable the development of pre-trained models that can jointly

Parameters

Temperature 0.3

Max length (tokens) 250

Stop sequences

Top probabilities 1

Frequency penalty 0

Presence penalty 0

Pre-response text

Post-response text

Learn more

Generate Undo Regenerate Tokens: 571

IBM watsonx.ai is a studio of integrated tools for working with generative AI capabilities that are powered by foundation models and for building machine learning models.

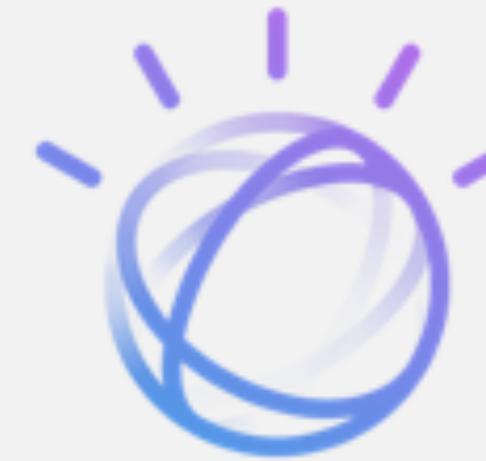
IBM watsonx.ai provides a secure and collaborative environment where you can access your organization's trusted data, automate AI processes, and deliver AI in your applications.



Class exercise

IBM watsonx studio

watsonx



Summarization



- Meeting transcript summary
- Earnings call summary

Classification



- Scenario classification
- Sentiment classification

Generation



- Marketing email generation
- Thank you note generation

Extraction



- Named entity extraction
- Fact extraction

Question and Answering



- Questions about an article
- Finance Q&A

Code



- Code generation
- Code translation

Class exercise

IBM watsonx studio

The screenshot shows the IBM WatsonX Studio interface. On the left, there's a sidebar with sections like 'Sample prompts', 'Summarization', 'Classification', 'Generation', 'Extraction', and 'Question answering'. The main area has tabs for 'Structured' and 'Freeform'. It includes fields for 'Instruction (optional)' and 'Examples (optional)'. A pink arrow points from the 'Instruction' field towards the model list on the right. The right side lists several models with their details:

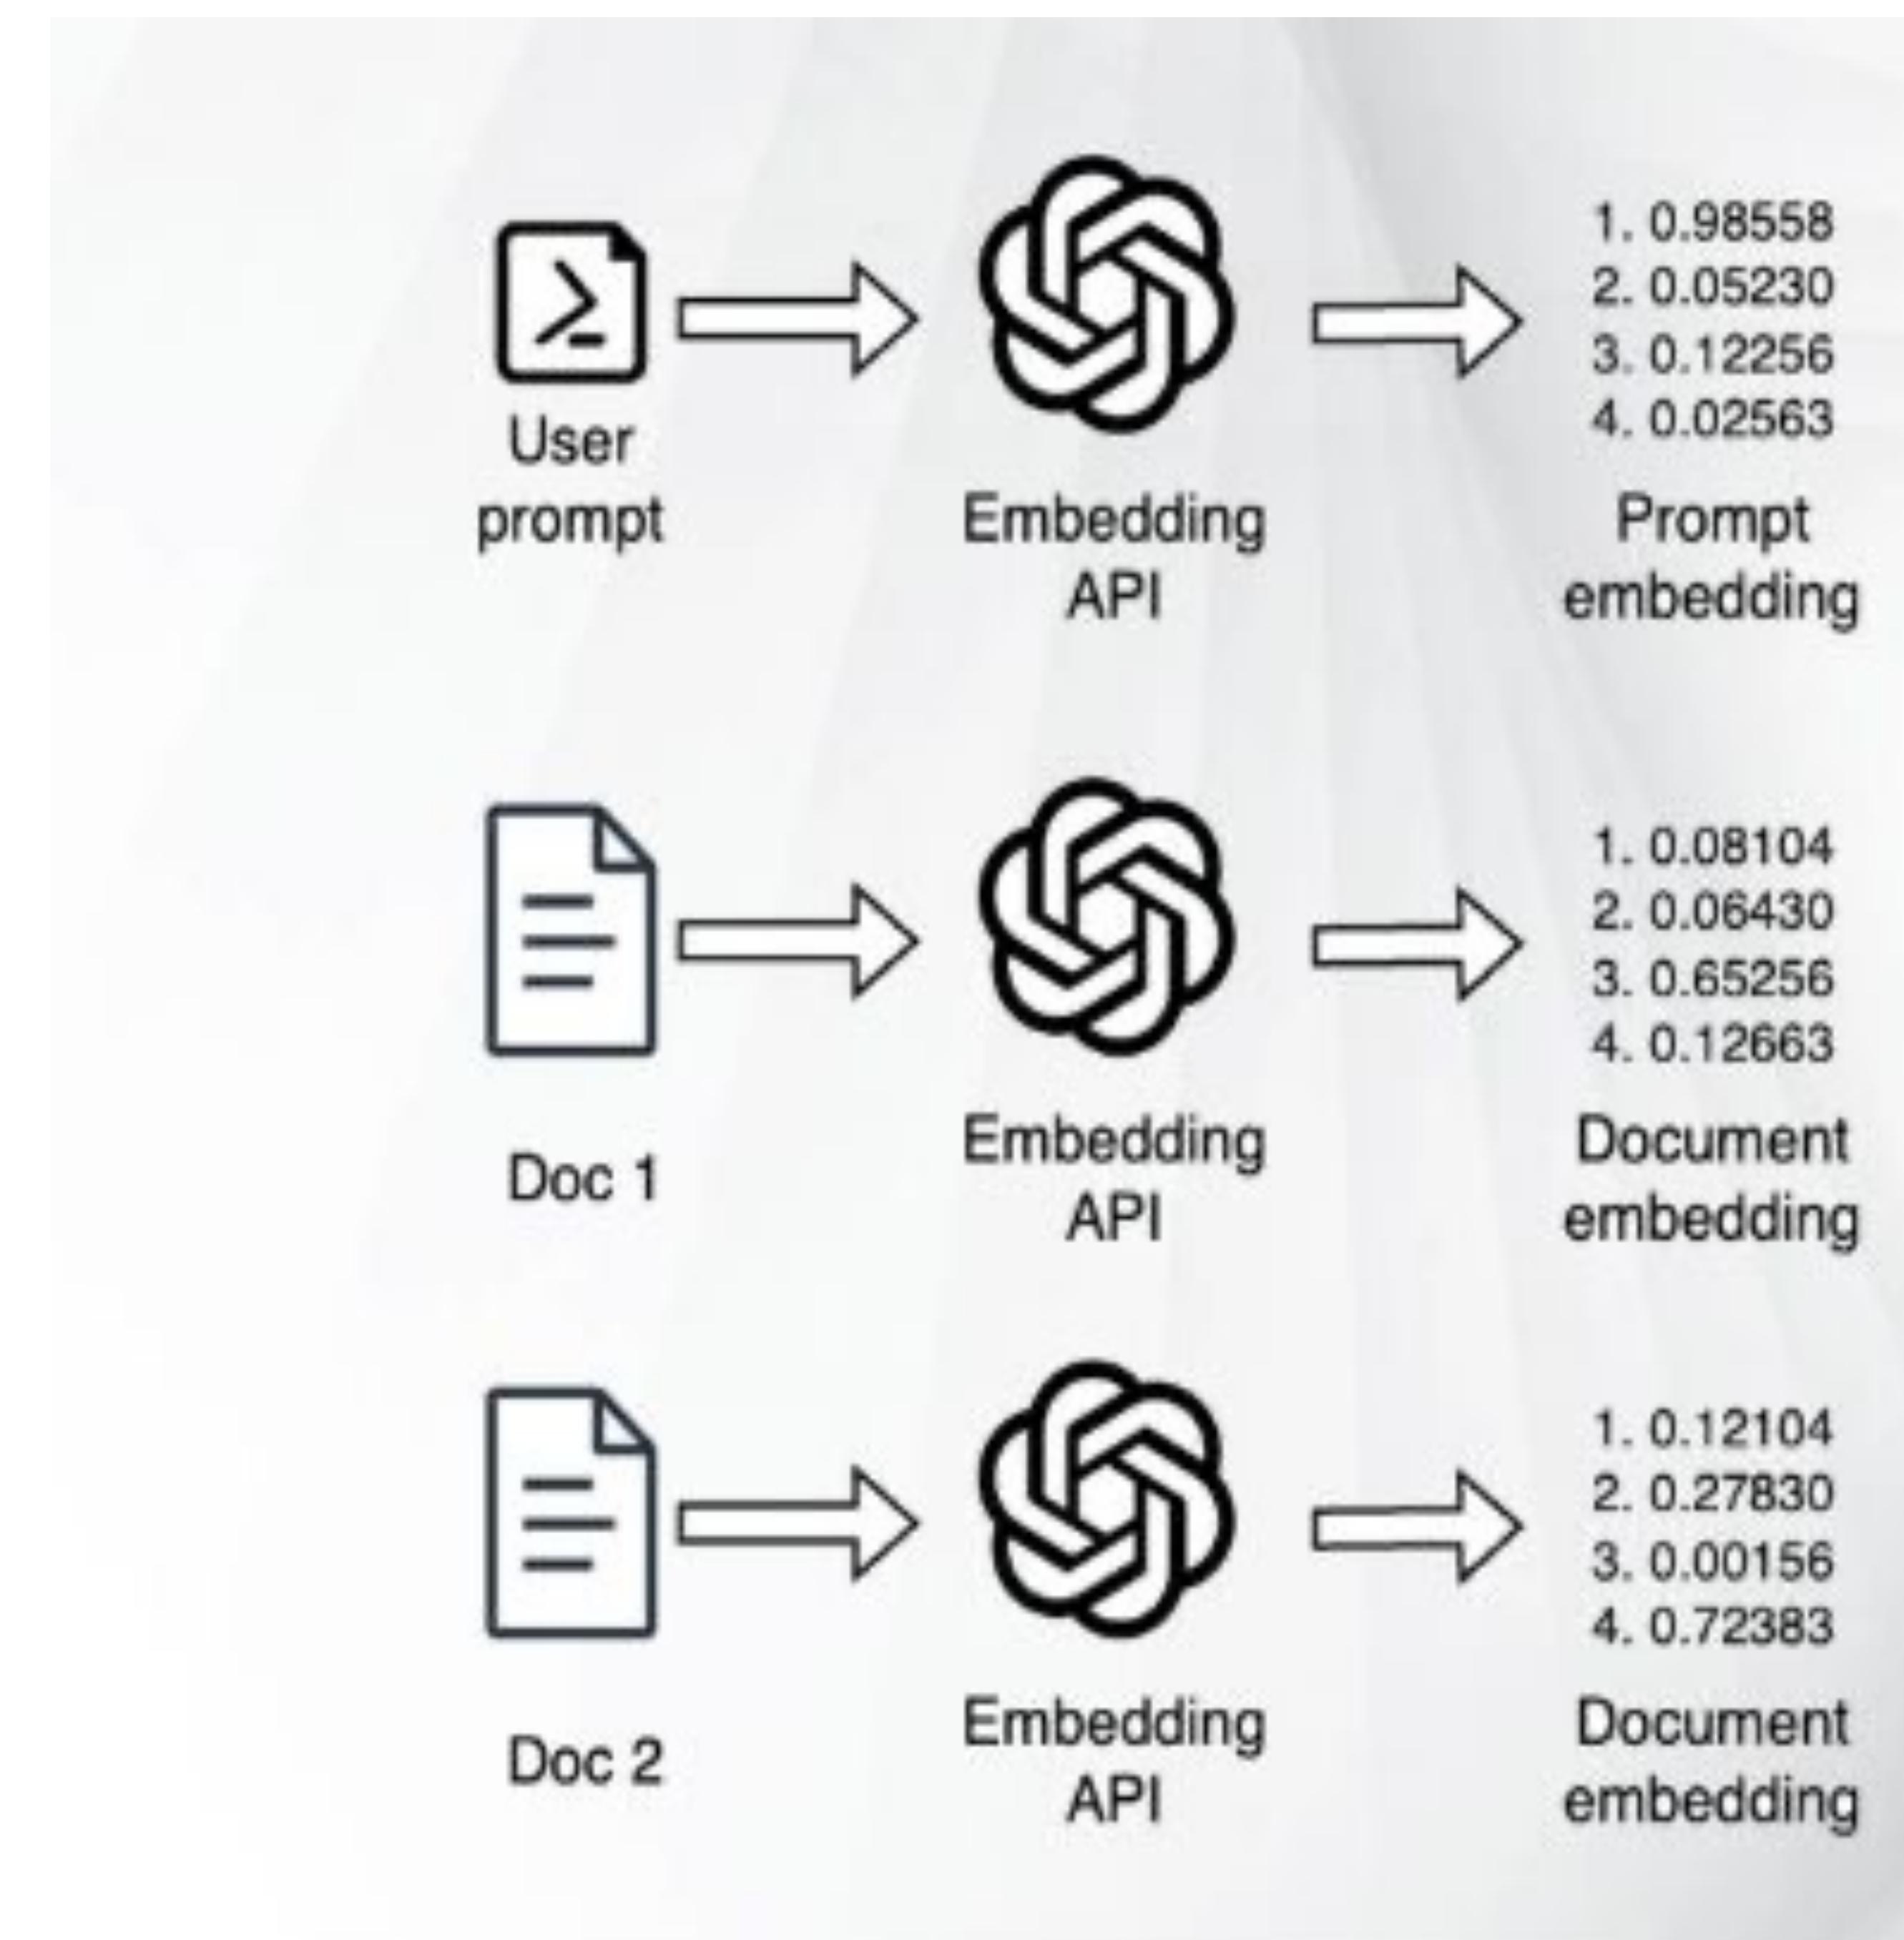
Model	Provider	Source
flan-ul2-20b	Google	Hugging Face
starcoder-15.5b	BigCode	Hugging Face
mt0-xxl-13b	BigScience	Hugging Face
gpt-neox-20b	EleutherAI	Hugging Face
flan-t5-xxl-11b	Google	Hugging Face
granite-13b-chat-v1	IBM	IBM
granite-13b-instruct-v1	IBM	IBM
mpt-7b-instruct2	Mosaic, tuned by IBM	Hugging Face
llama-2-70b-chat	Meta	Hugging Face

A 'Generate' button is located at the bottom center of the main workspace.

Azure OpenAI Studio

<https://platform.openai.com/account/api-keys>

Embeddings



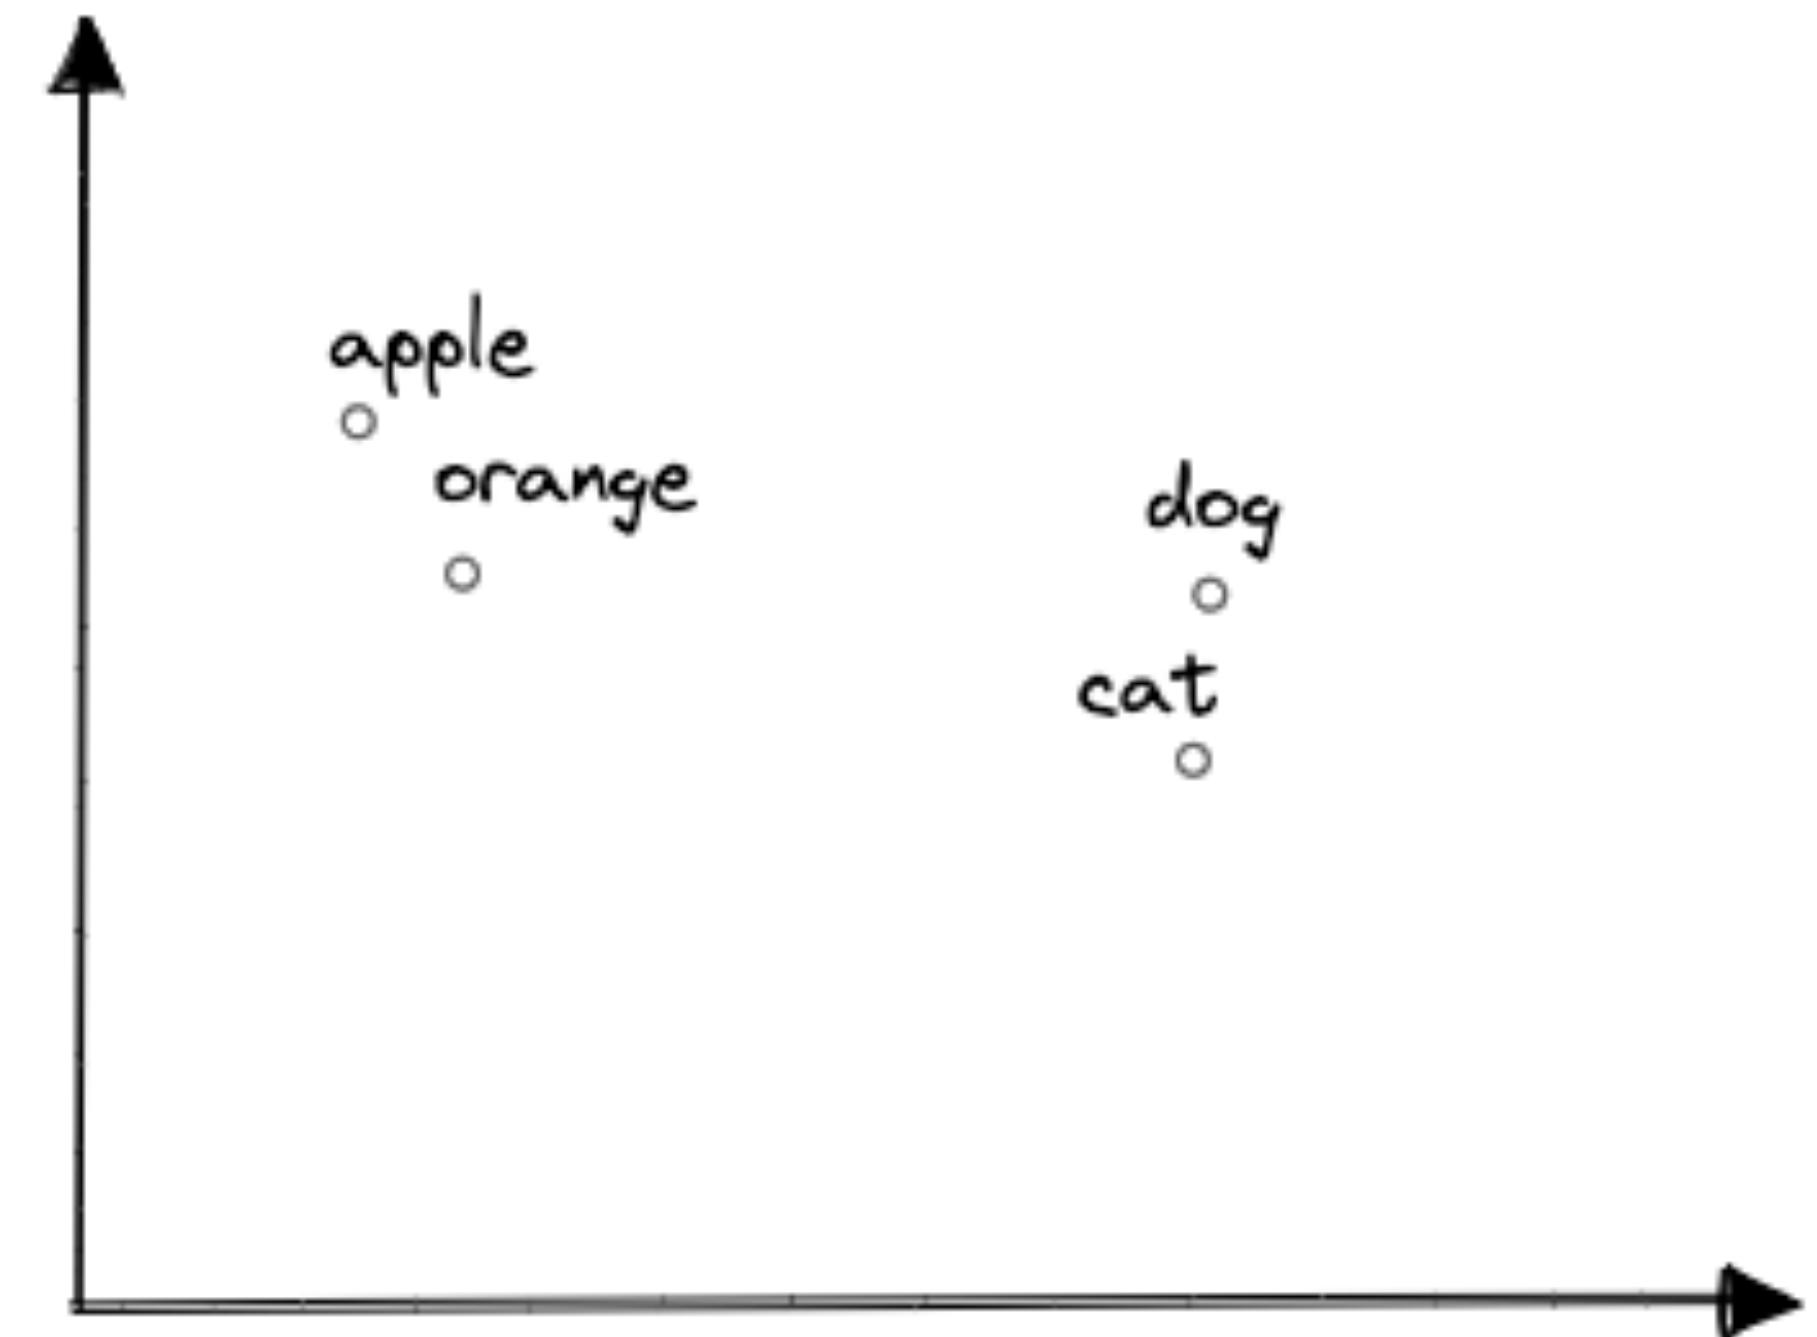
Embeddings

	dog	cat
0	0	
0	0	
0	1	
.	.	
.	.	
.	.	
1	0	
0	0	
0	0	
0	0	

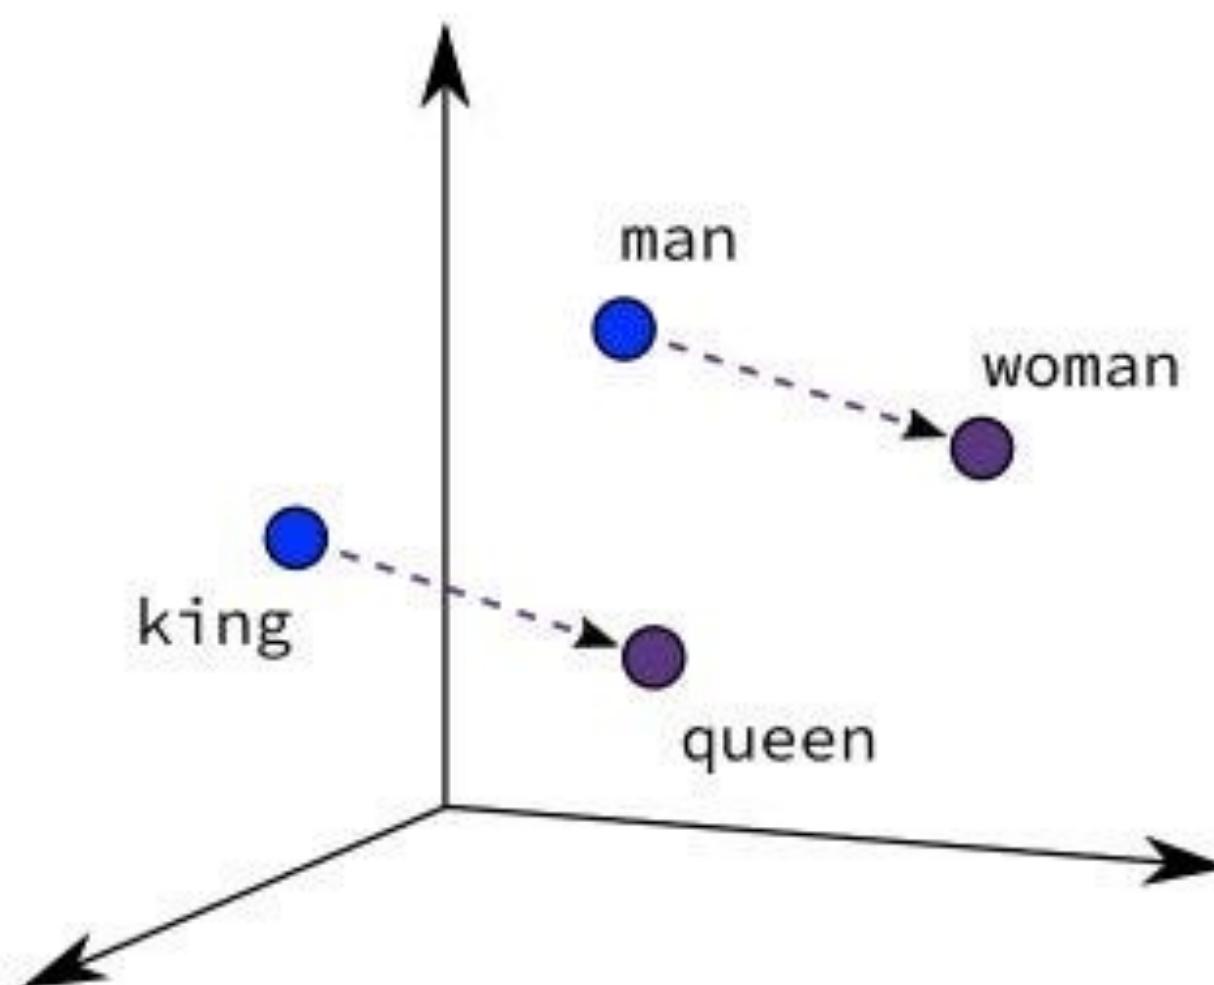


	human	fruit
0	.	
0	.	
0	.	
.	.	
.	.	
.	.	
1	0	
0	0	
0	0	
0	0	

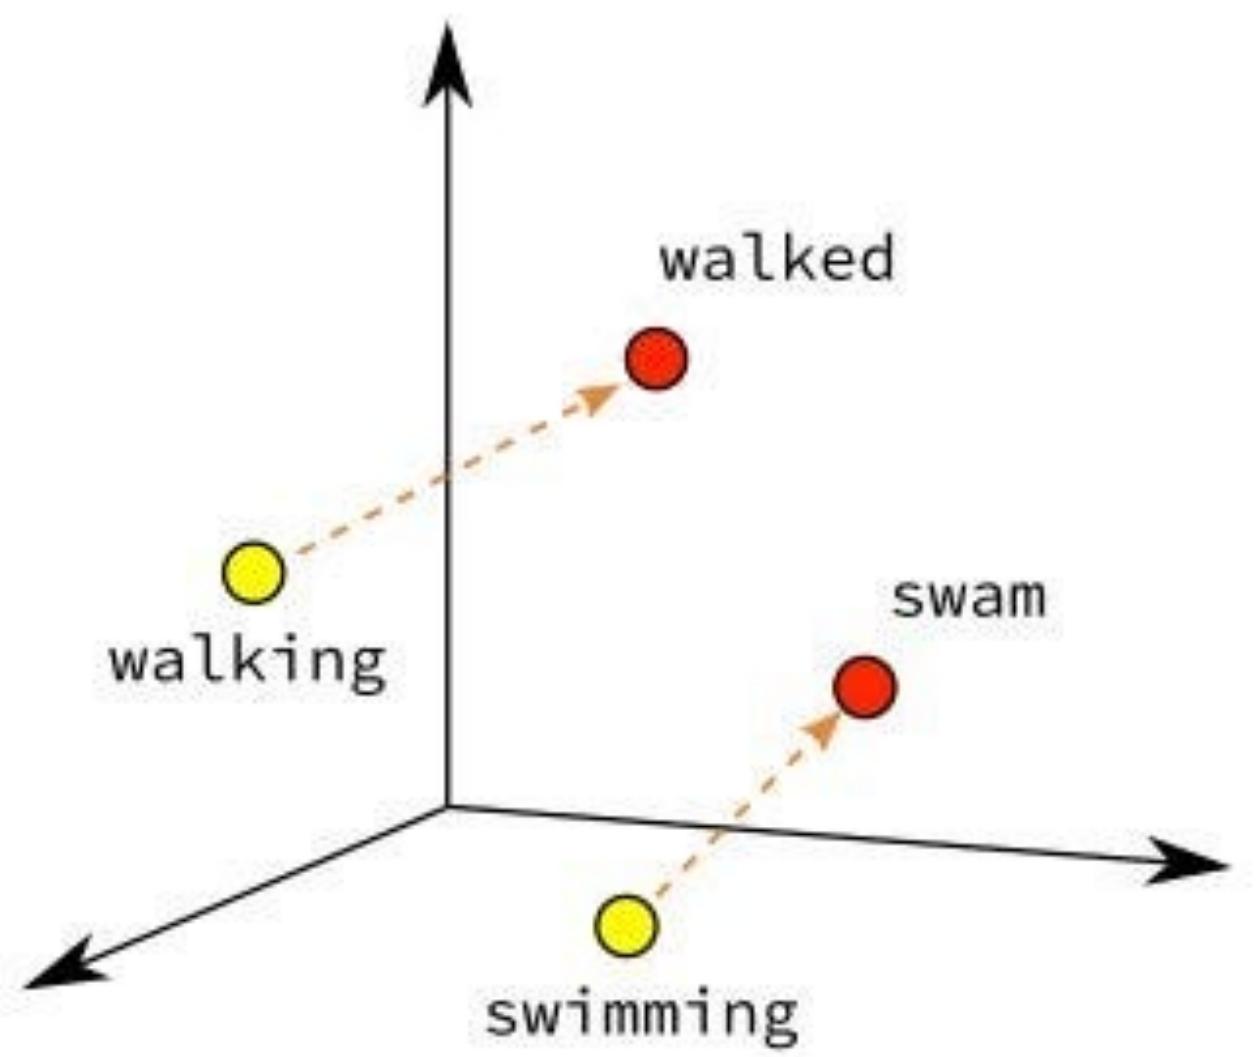
	dog	cat
0.05	0.01	
0.09	0.03	
.	.	
.	.	
.	.	
0.91	0.93	
0.08	0.04	



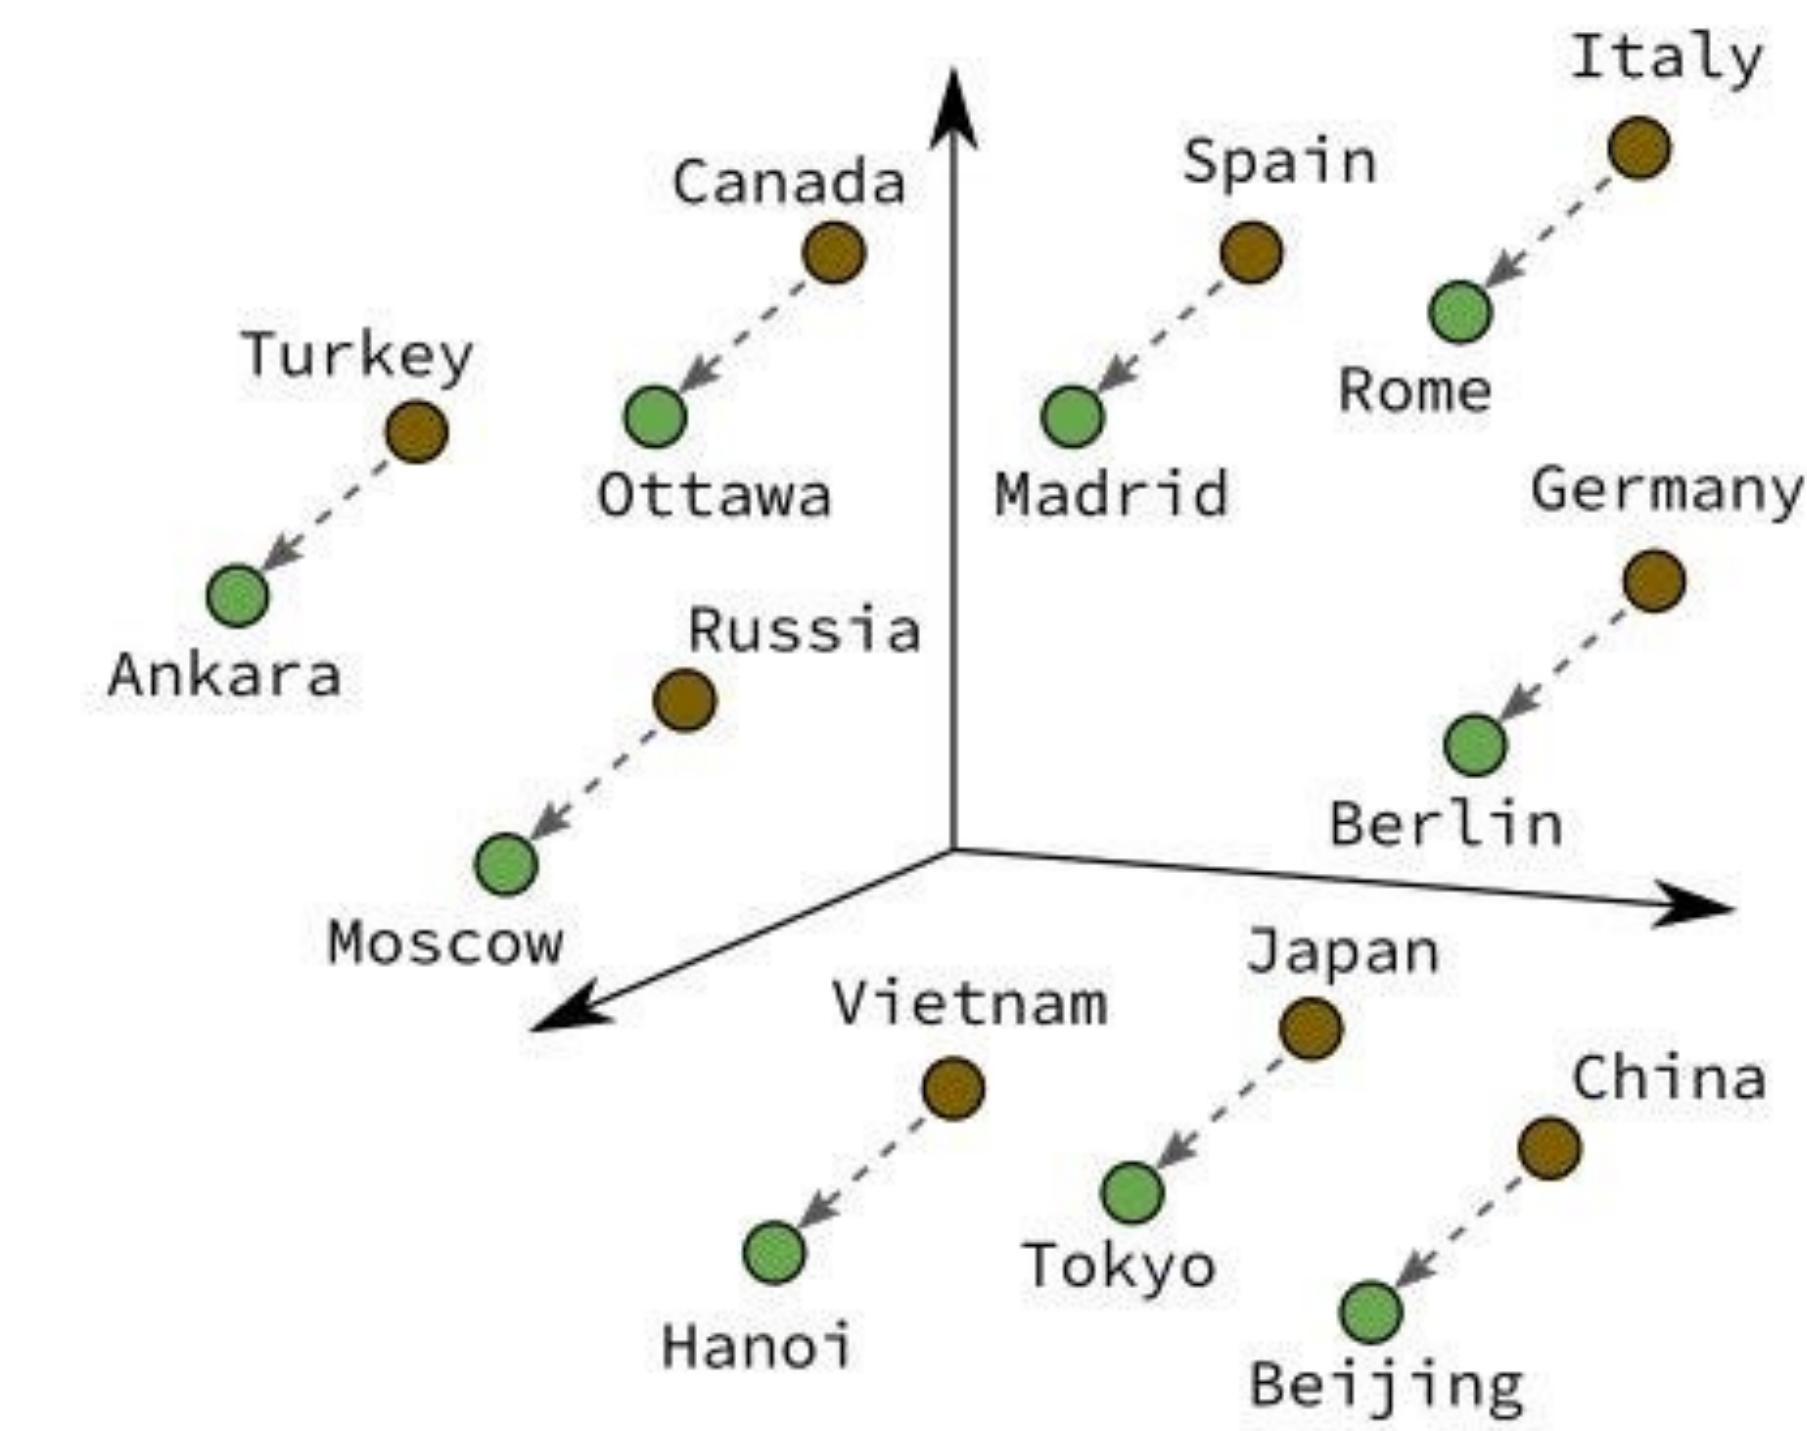
Embeddings



Male-Female

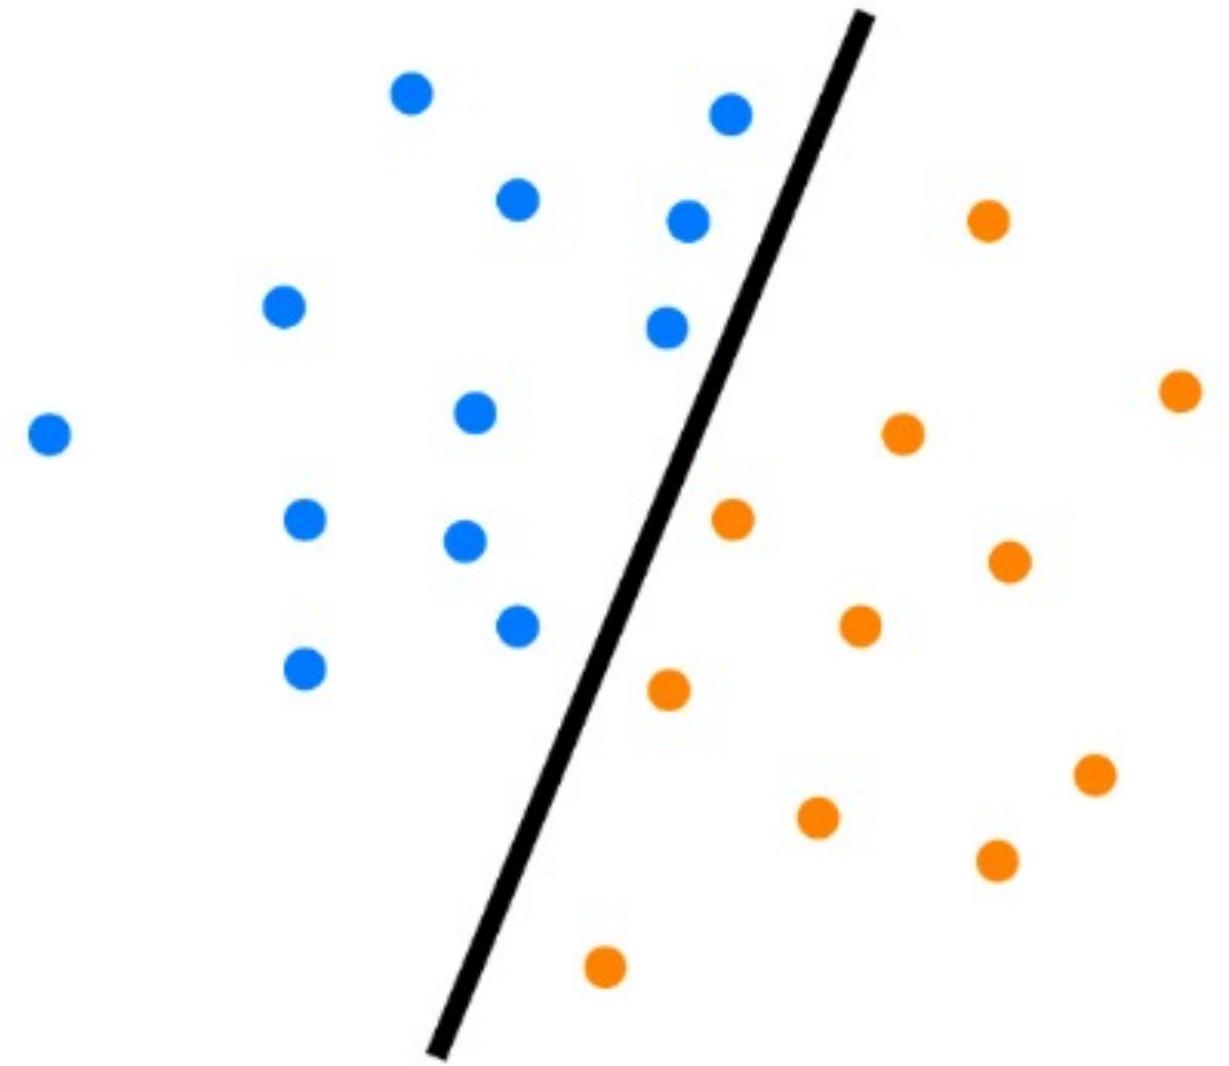


Verb Tense



Country-Capital

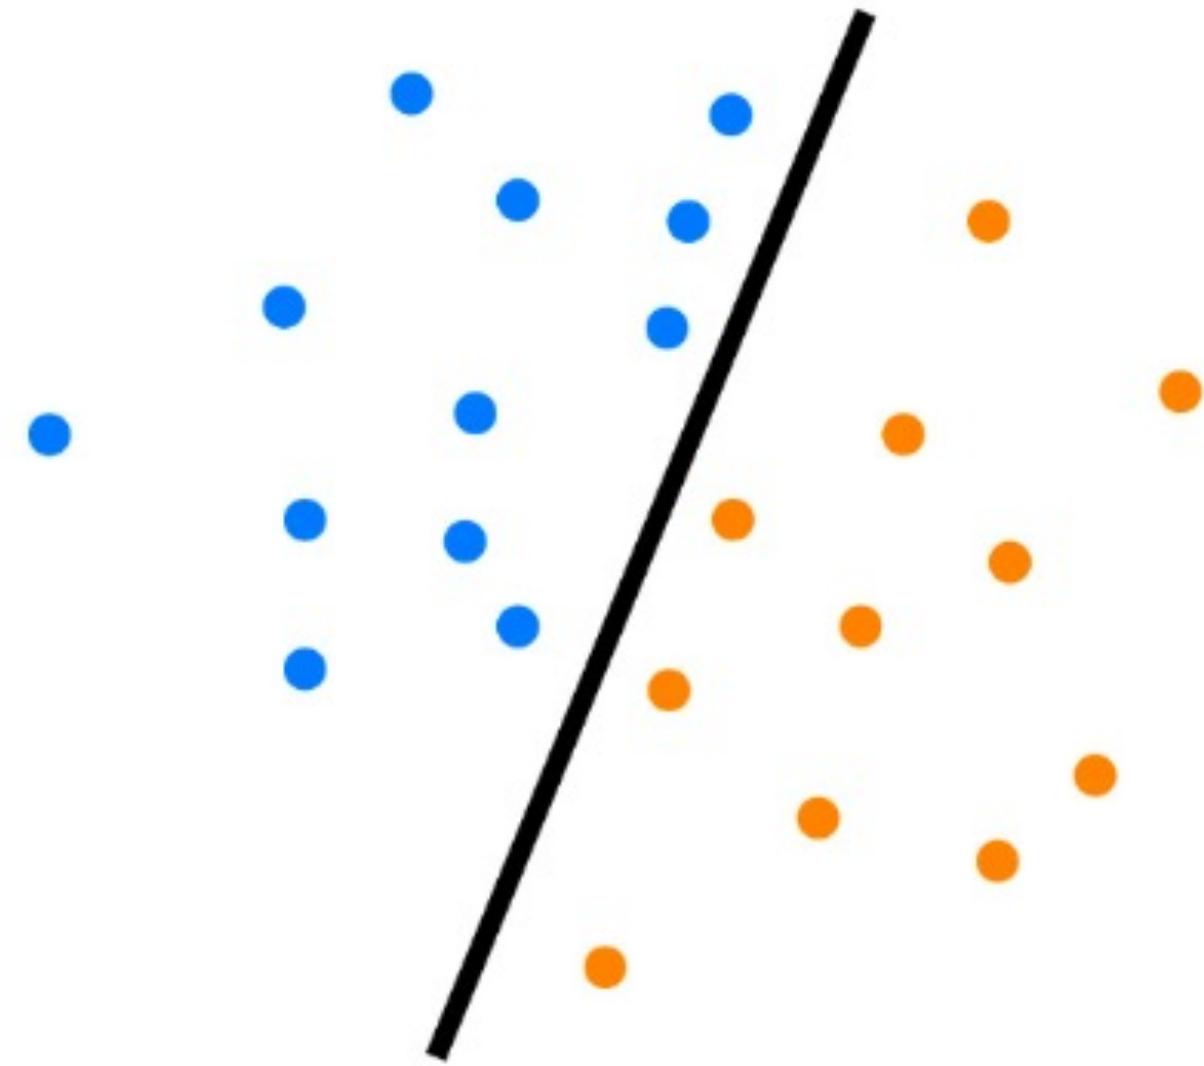
Embeddings For ...



- 1. Search**
- 2. Clustering**
- 3. Recommendations**
- 4. Anomaly detection**
- 5. Diversity measurement**
- 6. Classification**

1. where results are ranked by relevance to a query string
2. where text strings are grouped by similarity
3. where items with related text strings are recommended
4. where outliers with little relatedness are identified
5. where similarity distributions are analyzed
6. where text strings are classified by their most similar label

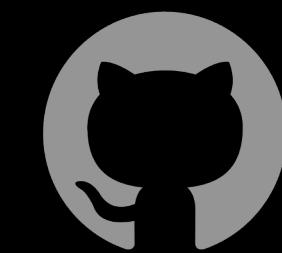
Embeddings For ...



- 1. Search**
- 2. Clustering**
- 3. Recommendations**
- 4. Anomaly detection**
- 5. Diversity measurement**
- 6. Classification**

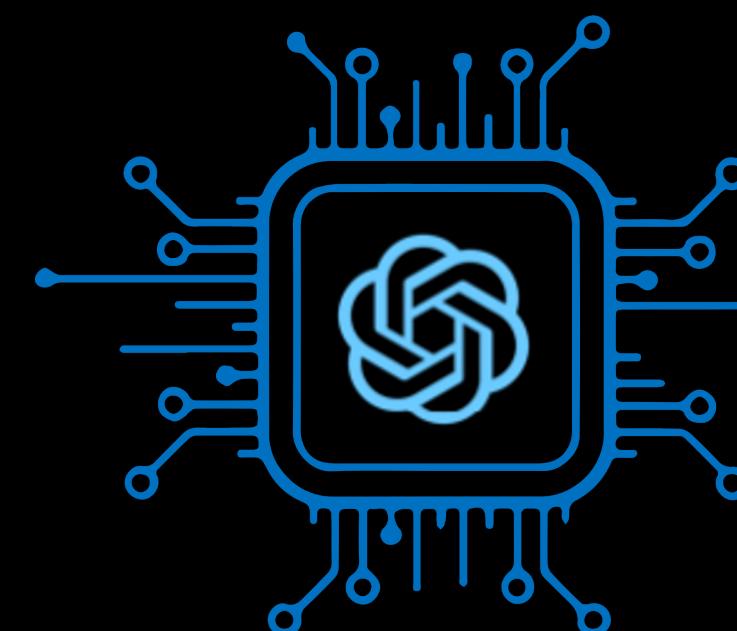
1. where results are ranked by relevance to a query string
2. where text strings are grouped by similarity
3. where items with related text strings are recommended
4. where outliers with little relatedness are identified
5. where similarity distributions are analyzed
6. where text strings are classified by their most similar label

Backup



OpenAI GitHub

OpenAI Codex
Model



Public code and text
on the internet

GitHub



GitHub
Copilot Service



Don't fly solo.

- Provide editor context
- Provide suggestions
- Improve suggestions

JS fetch_pic.js

push_to

```
1 const fetchNASAPictureOfTheDay = () => {
2   return fetch('https://api.nasa.gov/mars/rovers/curiosity/images/latest')
3     .then(response => response.json())
4     .then(json => {
5       return json;
6     });
7 }
```

Copilot



Azure AI

